

`fitr`

July 1, 2018

# Contents

<b>1</b>	<b>Overview &amp; Foundations</b>	<b>3</b>
<b>2</b>	<b>Tutorials</b>	<b>4</b>
	Getting Started . . . . .	4
	Installation . . . . .	4
	Simulating and Fitting a Two-Armed Bandit . . . . .	4
<b>I</b>	<b>API</b>	<b>6</b>
<b>3</b>	<b>Environments</b>	<b>7</b>
	<code>fitr.environments</code> . . . . .	7
	Graph . . . . .	7
	TwoArmedBandit . . . . .	11
	OrthogonalGoNoGo . . . . .	14
	TwoStep . . . . .	17
	ReverseTwoStep . . . . .	21
	RandomContextualBandit . . . . .	24
<b>4</b>	<b>Agents</b>	<b>28</b>
	<code>fitr.agents</code> . . . . .	28
	SoftmaxPolicy . . . . .	28
	StickySoftmaxPolicy . . . . .	29
	EpsilonGreedyPolicy . . . . .	31
	ValueFunction . . . . .	32
	DummyLearner . . . . .	34
	InstrumentalRescorlaWagnerLearner . . . . .	36
	QLearner . . . . .	39
	SARSA Learner . . . . .	42
	Agent . . . . .	44
	BanditAgent . . . . .	45
	MDP Agent . . . . .	47
	RandomBanditAgent . . . . .	48
	RandomMDP Agent . . . . .	50
	Notes . . . . .	50
	SARSA Softmax Agent . . . . .	51
	QLearning Softmax Agent . . . . .	53
	RW Softmax Agent . . . . .	55

	RWSoftmaxAgentRewardSensitivity . . . . .	57
<b>5</b>	<b>Data</b>	<b>60</b>
	fitr.data . . . . .	60
	BehaviouralData . . . . .	60
	merge_behavioural_data . . . . .	62
<b>6</b>	<b>Inference</b>	<b>63</b>
	fitr.inference . . . . .	63
	OptimizationResult . . . . .	63
	mlepar . . . . .	64
	l_bfgs_b . . . . .	64
<b>7</b>	<b>Criticism</b>	<b>66</b>
	fitr.criticism . . . . .	66
	actual_estimate . . . . .	66
<b>8</b>	<b>Metrics</b>	<b>67</b>
	fitr.metrics . . . . .	67
	bic . . . . .	67
	linear_correlation . . . . .	67
	lme . . . . .	68
	log_loss . . . . .	68
<b>9</b>	<b>Utilities</b>	<b>69</b>
	fitr.utils . . . . .	69
	logsumexp . . . . .	69
	relu . . . . .	69
	scale_data . . . . .	70
	sigmoid . . . . .	70
	softmax . . . . .	71
	stable_exp . . . . .	71

## **Chapter 1**

# **Overview & Foundations**

## Chapter 2

# Tutorials

### Getting Started

#### Installation

```
pip install git+https://github.com/abrahamnunes/fitr.git
```

### Simulating and Fitting a Two-Armed Bandit

```
import numpy as np
import matplotlib.pyplot as plt
from fitr import generate_behavioural_data
from fitr.environments import TwoArmedBandit
from fitr.agents import RWSoftmaxAgent
from fitr.inference import mlepar
from fitr.utils import sigmoid
from fitr.utils import relu
from fitr.criticism.plotting import actual_estimate

N = 50  # number of subjects
T = 200 # number of trials

# Generate synthetic data
data = generate_behavioural_data(TwoArmedBandit(), RWSoftmaxAgent, N, T)

# Create log-likelihood function
def log_prob(w, D):
    lr = sigmoid(w[0], a_min=-6, a_max=6)
    ist = relu(w[1], a_max=10)
    agent = RWSoftmaxAgent(TwoArmedBandit(), lr, ist)
    L = 0
    for t in range(D.shape[0]):
        x = D[t, :3]
```

```
    u = D[t, 3:5]
    r = D[t, 5]
    x_ = D[t, 6:]
    L += u@agent.log_prob(x)
    agent.learning(x, u, r, x_, None)
    return L

# Fit model
res = mlepar(log_prob, data.tensor, nparams=2, maxstarts=5)
X = res.transform_xmin([sigmoid, relu])

# Criticism
lr_fig = actual_estimate(data.params[:,1], X[:,0]); plt.show()
ist_fig = actual_estimate(data.params[:,2], X[:,1]); plt.show()
```

# **Part I**

# **API**

## Chapter 3

# Environments

### `fitr.environments`

Functions to synthesize data from behavioural tasks.

### Graph

`fitr.environments.Graph()`

Base object that defines a reinforcement learning task.

### Definitions

- $\mathbf{x} \in \mathcal{X}$  be a one-hot state vector, where  $|\mathcal{X}| = n_x$
- $\mathbf{u} \in \mathcal{U}$  be a one-hot action vector, where  $|\mathcal{U}| = n_u$
- $\mathbf{T} = p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$  be a transition tensor
- $p(\mathbf{x})$  be a distribution over starting states
- $\mathcal{J} : \mathcal{X} \rightarrow \mathcal{R}$ , where  $\mathcal{R} \subseteq \mathbb{R}$  be a reward function

Arguments:

- **T**: Transition tensor
- **R**: Vector of rewards for each state such that scalar reward  $r_t = \mathbf{r}^o p \mathbf{x}$
- **end\_states**: A vector  $\{0, 1\}^{n_x}$  identifying which states terminate a trial (aka episode)
- **p\_start**: Initial state distribution
- **label**: A string identifying a name for the task
- **state\_labels**: A list or array of strings labeling the different states (for plotting purposes)
- **action\_labels**: A list or array of strings labeling the different actions (for plotting purposes)
- **rng**: `np.random.RandomState` object
- **f\_reward**: A function whose first argument is a vector of rewards for each state, and whose second argument is a state vector, and whose output is a scalar reward
- **cmap**: Matplotlib colormap for plotting.



**Notes**

There are two critical methods for the Graph class: `observation()` and `step`. All instances of a Graph must be able to call these functions. Let's say you have some bandit task `MyBanditTask` that inherits from Graph. To run such a task would look something like this:

```
env = MyBanditTask()           # Instantiate your environment object
agent = MyAgent()              # Some agent object (arbitrary, really)
for t in range(ntrials):
    x = env.observation()       # Samples initial state
    u = agent.action(x)         # Choose some action
    x_, r, done = agent.step(u) # Transition based on action
```

What differentiates tasks are the transition tensor  $T$ , starting state distribution  $p(\mathbf{x})$  and reward function  $\mathcal{J}$  (which here would include the reward vector  $\mathbf{r}$ ).

**Graph.adjacency\_matrix\_decomposition**

```
fitr.environments.adjacency_matrix_decomposition(self)
```

Singular value decomposition of the graph adjacency matrix

**Graph.get\_graph\_depth**

```
fitr.environments.get_graph_depth(self)
```

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START`->node & node->`END` transitions.

Returns:

An int identifying the depth of the current graph for a single trial of the task

**Graph.laplacian\_matrix\_decomposition**

```
fitr.environments.laplacian_matrix_decomposition(self)
```

Singular value decomposition of the graph Laplacian

**Graph.make\_action\_labels**

```
fitr.environments.make_action_labels(self)
```

Creates labels for the actions (for plotting) if none provided

---

### **Graph.make\_digraph**

```
fitr.environments.make_digraph(self)
```

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

### **Graph.make\_state\_labels**

```
fitr.environments.make_state_labels(self)
```

Creates labels for the states (for plotting) if none provided

---

### **Graph.make\_undirected\_graph**

```
fitr.environments.make_undirected_graph(self)
```

Converts the `DiGraph` to undirected and computes some stats

---

### **Graph.observation**

```
fitr.environments.observation(self)
```

Samples an initial state from the start-state distribution  $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

**Graph.plot\_action\_outcome\_probabilities**

```
fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None)
```

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

**Graph.plot\_graph**

```
fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.)
```

Plots the directed graph of the task

**Graph.plot\_spectral\_properties**

```
fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile)
```

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

**Graph.random\_action**

```
fitr.environments.random_action(self)
```

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\left(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\right)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

**Graph.step**

```
fitr.environments.step(self, action)
```

Executes a state transition in the environment.

Arguments:

action : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

**TwoArmedBandit**

```
fitr.environments.TwoArmedBandit()
```

Two armed bandit just as a tester

---

**TwoArmedBandit.adjacency\_matrix\_decomposition**

```
fitr.environments.adjacency_matrix_decomposition(self)
```

Singular value decomposition of the graph adjacency matrix

---

**TwoArmedBandit.get\_graph\_depth**

```
fitr.environments.get_graph_depth(self)
```

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**TwoArmedBandit.laplacian\_matrix\_decomposition**

```
fitr.environments.laplacian_matrix_decomposition(self)
```

Singular value decomposition of the graph Laplacian

---

**TwoArmedBandit.make\_action\_labels**

```
fitr.environments.make_action_labels(self)
```

Creates labels for the actions (for plotting) if none provided

---

**TwoArmedBandit.make\_digraph**

```
fitr.environments.make_digraph(self)
```

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

**TwoArmedBandit.make\_state\_labels**

```
fitr.environments.make_state_labels(self)
```

Creates labels for the states (for plotting) if none provided

---

**TwoArmedBandit.make\_undirected\_graph**

```
fitr.environments.make_undirected_graph(self)
```

Converts the `DiGraph` to undirected and computes some stats

---

**TwoArmedBandit.observation**

```
fitr.environments.observation(self)
```

Samples an initial state from the start-state distribution  $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

### **TwoArmedBandit.plot\_action\_outcome\_probabilities**

```
fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None)
```

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

### **TwoArmedBandit.plot\_graph**

```
fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.)
```

Plots the directed graph of the task

---

### **TwoArmedBandit.plot\_spectral\_properties**

```
fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile)
```

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

### **TwoArmedBandit.random\_action**

```
fitr.environments.random_action(self)
```

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\left(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\right)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

**TwoArmedBandit.step**

```
fitr.environments.step(self, action)
```

Executes a state transition in the environment.

Arguments:

`action` : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

**OrthogonalGoNoGo**

```
fitr.environments.OrthogonalGoNoGo()
```

The orthogonal GoNogo task from Guitart-Masip et al. (2012)

---

**OrthogonalGoNoGo.adjacency\_matrix\_decomposition**

```
fitr.environments.adjacency_matrix_decomposition(self)
```

Singular value decomposition of the graph adjacency matrix

---

**OrthogonalGoNoGo.get\_graph\_depth**

```
fitr.environments.get_graph_depth(self)
```

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**OrthogonalGoNoGo.laplacian\_matrix\_decomposition**

```
fitr.environments.laplacian_matrix_decomposition(self)
```

Singular value decomposition of the graph Laplacian

---

**OrthogonalGoNoGo.make\_action\_labels**

```
fitr.environments.make_action_labels(self)
```

Creates labels for the actions (for plotting) if none provided

---

**OrthogonalGoNoGo.make\_digraph**

```
fitr.environments.make_digraph(self)
```

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

**OrthogonalGoNoGo.make\_state\_labels**

```
fitr.environments.make_state_labels(self)
```

Creates labels for the states (for plotting) if none provided

---

**OrthogonalGoNoGo.make\_undirected\_graph**

```
fitr.environments.make_undirected_graph(self)
```

Converts the `DiGraph` to undirected and computes some stats

---

**OrthogonalGoNoGo.observation**

```
fitr.environments.observation(self)
```

Samples an initial state from the start-state distribution  $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:



A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

### **OrthogonalGoNoGo.plot\_action\_outcome\_probabilities**

```
fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None)
```

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

### **OrthogonalGoNoGo.plot\_graph**

```
fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.)
```

Plots the directed graph of the task

---

### **OrthogonalGoNoGo.plot\_spectral\_properties**

```
fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile)
```

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

### **OrthogonalGoNoGo.random\_action**

```
fitr.environments.random_action(self)
```

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\left(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\right)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

### OrthogonalGoNoGo.step

```
fitr.environments.step(self, action)
```

Executes a state transition in the environment.

Arguments:

`action` : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

### TwoStep

```
fitr.environments.TwoStep()
```

An implementation of the Two-Step Task from Daw et al. (2011).

Arguments:

- **mu**: float identifying the drift of the reward-determining Gaussian random walks
  - **sd**: float identifying the standard deviation of the reward-determining Gaussian random walks
- 

### TwoStep.adjacency\_matrix\_decomposition

```
fitr.environments.adjacency_matrix_decomposition(self)
```

Singular value decomposition of the graph adjacency matrix

---

### TwoStep.f\_reward

```
fitr.environments.f_reward(self, R, x)
```

---

**TwoStep.get\_graph\_depth**

```
fitr.environments.get_graph_depth(self)
```

Returns the depth of the task graph.

Calculated as the depth from *START* (pre-initial state) to *END* (which absorbs trial from all terminal states), minus 2 to account for the *START*->node & node->*END* transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**TwoStep.laplacian\_matrix\_decomposition**

```
fitr.environments.laplacian_matrix_decomposition(self)
```

Singular value decomposition of the graph Laplacian

---

**TwoStep.make\_action\_labels**

```
fitr.environments.make_action_labels(self)
```

Creates labels for the actions (for plotting) if none provided

---

**TwoStep.make\_digraph**

```
fitr.environments.make_digraph(self)
```

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

**TwoStep.make\_state\_labels**

```
fitr.environments.make_state_labels(self)
```

Creates labels for the states (for plotting) if none provided

---

**TwoStep.make\_undirected\_graph**

```
fitr.environments.make_undirected_graph(self)
```

Converts the DiGraph to undirected and computes some stats

---

**TwoStep.observation**

```
fitr.environments.observation(self)
```

Samples an initial state from the start-state distribution  $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

**TwoStep.plot\_action\_outcome\_probabilities**

```
fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None)
```

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

**TwoStep.plot\_graph**

```
fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.)
```

Plots the directed graph of the task

---

**TwoStep.plot\_reward\_paths**

```
fitr.environments.plot_reward_paths(self, outfile=None, outfiletype='pdf', figsize=N
```

---

**TwoStep.plot\_spectral\_properties**

```
fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile)
```

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

**TwoStep.random\_action**

```
fitr.environments.random_action(self)
```

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\left(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\right)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

**TwoStep.step**

```
fitr.environments.step(self, action)
```

Executes a state transition in the environment.

Arguments:

`action` : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

**ReverseTwoStep**

```
fitr.environments.ReverseTwoStep()
```

From Kool & Gershman 2016.

---

**ReverseTwoStep.adjacency\_matrix\_decomposition**

```
fitr.environments.adjacency_matrix_decomposition(self)
```

Singular value decomposition of the graph adjacency matrix

---

**ReverseTwoStep.f\_reward**

```
fitr.environments.f_reward(self, R, x)
```

---

**ReverseTwoStep.get\_graph\_depth**

```
fitr.environments.get_graph_depth(self)
```

Returns the depth of the task graph.

Calculated as the depth from *START* (pre-initial state) to *END* (which absorbs trial from all terminal states), minus 2 to account for the *START*->node & node->*END* transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**ReverseTwoStep.laplacian\_matrix\_decomposition**

```
fitr.environments.laplacian_matrix_decomposition(self)
```

Singular value decomposition of the graph Laplacian

---

**ReverseTwoStep.make\_action\_labels**

```
fitr.environments.make_action_labels(self)
```

Creates labels for the actions (for plotting) if none provided

---

**ReverseTwoStep.make\_digraph**

```
fitr.environments.make_digraph(self)
```

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

**ReverseTwoStep.make\_state\_labels**

```
fitr.environments.make_state_labels(self)
```

Creates labels for the states (for plotting) if none provided

---

**ReverseTwoStep.make\_undirected\_graph**

```
fitr.environments.make_undirected_graph(self)
```

Converts the `DiGraph` to undirected and computes some stats

---

**ReverseTwoStep.observation**

```
fitr.environments.observation(self)
```

Samples an initial state from the start-state distribution  $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

**ReverseTwoStep.plot\_action\_outcome\_probabilities**

```
fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None)
```

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

**ReverseTwoStep.plot\_graph**

```
fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.
```

Plots the directed graph of the task

---

**ReverseTwoStep.plot\_spectral\_properties**

```
fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile
```

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

**ReverseTwoStep.random\_action**

```
fitr.environments.random_action(self)
```

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\left(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\right)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

**ReverseTwoStep.step**

```
fitr.environments.step(self, action)
```

Executes a state transition in the environment.

Arguments:

`action` : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.



---

## RandomContextualBandit

```
fitr.environments.RandomContextualBandit()
```

Generates a random bandit task

Arguments:

- **nactions**: Number of actions
  - **noutcomes**: Number of outcomes
  - **nstates**: Number of contexts
  - **min\_actions\_per\_context**: Different contexts may have more or fewer actions than others (never more than `nactions`). This variable describes the minimum number of actions allowed in a context.
  - **alpha**:
  - **alpha\_start**:
  - **shift\_flip**:
  - **reward\_lb**: Lower bound for drifting rewards
  - **reward\_ub**: Upper bound for drifting rewards
  - **reward\_drift**: Values (on or off) determining whether rewards are allowed to drift
  - **drift\_mu**: Mean of the Gaussian random walk determining reward
  - **drift\_sd**: Standard deviation of Gaussian random walk determining reward
- 

## RandomContextualBandit.adjacency\_matrix\_decomposition

```
fitr.environments.adjacency_matrix_decomposition(self)
```

Singular value decomposition of the graph adjacency matrix

---

## RandomContextualBandit.f\_reward

```
fitr.environments.f_reward(self, R, x)
```

---

## RandomContextualBandit.get\_graph\_depth

```
fitr.environments.get_graph_depth(self)
```

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START`->node & node->`END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**RandomContextualBandit.laplacian\_matrix\_decomposition**

```
fitr.environments.laplacian_matrix_decomposition(self)
```

Singular value decomposition of the graph Laplacian

---

**RandomContextualBandit.make\_action\_labels**

```
fitr.environments.make_action_labels(self)
```

Creates labels for the actions (for plotting) if none provided

---

**RandomContextualBandit.make\_digraph**

```
fitr.environments.make_digraph(self)
```

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

**RandomContextualBandit.make\_state\_labels**

```
fitr.environments.make_state_labels(self)
```

Creates labels for the states (for plotting) if none provided

---

**RandomContextualBandit.make\_undirected\_graph**

```
fitr.environments.make_undirected_graph(self)
```

Converts the `DiGraph` to undirected and computes some stats

---

**RandomContextualBandit.observation**

```
fitr.environments.observation(self)
```

Samples an initial state from the start-state distribution  $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray( (nstates, ) )` indicating the starting state.

Examples:

```
x = env.observation()
```

---

### **RandomContextualBandit.plot\_action\_outcome\_probabilities**

```
fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None)
```

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

### **RandomContextualBandit.plot\_graph**

```
fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.)
```

Plots the directed graph of the task

---

### **RandomContextualBandit.plot\_spectral\_properties**

```
fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile)
```

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

### **RandomContextualBandit.random\_action**

```
fitr.environments.random_action(self)
```

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\left(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\right)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

### **RandomContextualBandit.step**

```
fitr.environments.step(self, action)
```

Executes a state transition in the environment.

Arguments:

`action` : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

# Chapter 4

## Agents

### `fitr.agents`

A modular way to build and test reinforcement learning agents.

There are three main submodules:

- `fitr.agents.policies`: which describe a class of functions essentially representing  $f : \mathcal{X} \rightarrow \mathcal{U}$
- `fitr.agents.value_functions`: which describe a class of functions essentially representing  $\mathcal{V} : \mathcal{X} \rightarrow \mathbb{R}$  and/or  $\mathcal{Q} : \mathcal{Q} \times \mathcal{U} \rightarrow \mathbb{R}$
- `fitr.agents.agents`: classes of agents that are combinations of policies and value functions, along with some convenience functions for generating data from `fitr.environments.Graph environments`.

### SoftmaxPolicy

`fitr.agents.policies.SoftmaxPolicy()`

Action selection by sampling from a multinomial whose parameters are given by a softmax.

Action sampling is

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \zeta(\mathbf{v})).$$

Parameters of that distribution are

$$p(\mathbf{u}|\mathbf{v}) = \zeta(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i e^{\beta v_i}}.$$

Arguments:

- **inverse\_softmax\_temp**: Inverse softmax temperature  $\beta$
- **rng**: `np.random.RandomState` object

**SoftmaxPolicy.action\_prob**

```
fitr.agents.policies.action_prob(self, x)
```

Computes the softmax

---

**SoftmaxPolicy.log\_prob**

```
fitr.agents.policies.log_prob(self, x)
```

Computes the log-probability of an action **u**

$$\log p(\mathbf{u}|\mathbf{v}) = \beta \mathbf{v} - \log \sum_{v_i} e^{\beta v_i}$$

Arguments:

- **x**: State vector of type `ndarray( (nstates, ) )`

Returns:

Scalar log-probability

---

**SoftmaxPolicy.sample**

```
fitr.agents.policies.sample(self, x)
```

Samples from the action distribution

---

**StickySoftmaxPolicy**

```
fitr.agents.policies.StickySoftmaxPolicy()
```

Action selection by sampling from a multinomial whose parameters are given by a softmax, but with accounting for the tendency to persevere (i.e. choosing the previously used action without considering its value).

Let  $\mathbf{u}_{t-1} = (u_{t-1}^{(i)})_{i=1}^{|\mathcal{U}|}$  be a one hot vector representing the action taken at the last step, and  $\beta^\rho$  be an inverse softmax temperature for the influence of this last action.

Action sampling is thus:

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v}, \mathbf{u}_{t-1})).$$

Parameters of that distribution are

$$p(\mathbf{u}|\mathbf{v}, \mathbf{u}_{t-1}) = \varsigma(\mathbf{v}, \mathbf{u}_{t-1}) = \frac{e^{\beta\mathbf{v} + \beta^\rho \mathbf{u}_{t-1}}}{\sum_i e^{\beta v_i + \beta^\rho u_{t-1}^{(i)}}}.$$

Arguments:

- **inverse\_softmax\_temp**: Inverse softmax temperature  $\beta$
  - **perseveration**: Inverse softmax temperature  $\beta^\rho$  capturing the tendency to repeat the last action taken.
  - **rng**: `np.random.RandomState` object
- 

### StickySoftmaxPolicy.action\_prob

```
fitr.agents.policies.action_prob(self, x)
```

Computes the softmax

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` vector of action probabilities

---

### StickySoftmaxPolicy.log\_prob

```
fitr.agents.policies.log_prob(self, x)
```

Computes the log-probability of an action  $\mathbf{u}$

$$\log p(\mathbf{u}|\mathbf{v}, \mathbf{u}_{t-1}) = (\beta\mathbf{v} + \beta^\rho \mathbf{u}_{t-1}) - \log \sum_{v_i} e^{\beta v_i + \beta^\rho u_{t-1}^{(i)}}$$

Arguments:

- **x**: State vector of type `ndarray((nstates,))`

Returns:

Scalar log-probability

---

### StickySoftmaxPolicy.sample

```
fitr.agents.policies.sample(self, x)
```

Samples from the action distribution

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` one-hot action vector

---

## **EpsilonGreedyPolicy**

`fitr.agents.policies.EpsilonGreedyPolicy()`

A policy that takes the maximally valued action with probability  $1 - \epsilon$ , otherwise chooses randomlyself.

Arguments:

- **epsilon**: Probability of not taking the action with highest value
  - **rng**: `numpy.random.RandomState` object
- 

## **EpsilonGreedyPolicy.action\_prob**

`fitr.agents.policies.action_prob(self, x)`

Creates vector of action probabilities for e-greedy policy

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` vector of action probabilities

---

## **EpsilonGreedyPolicy.sample**

`fitr.agents.policies.sample(self, x)`

Samples from the action distribution

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` one-hot action vector

---



## ValueFunction

```
fitr.agents.value_functions.ValueFunction()
```

A general value function object.

A value function here is task specific and consists of several attributes:

- `nstates`: The number of states in the task,  $|\mathcal{X}|$
- `nactions`: Number of actions in the task,  $|\mathcal{U}|$
- `V`: State value function  $\mathbf{v} = \mathcal{V}(\mathbf{x})$
- `Q`: State-action value function  $\mathbf{Q} = \mathcal{Q}(\mathbf{x}, \mathbf{u})$
- `etrace`: An eligibility trace (optional)

Note that in general we rely on matrix-vector notation for value functions, rather than function notation. Vectors in the mathematical typesetting are by default column vectors.

Arguments:

- `env`: A `fitr.environments.Graph`
- 

## ValueFunction.Qmax

```
fitr.agents.value_functions.Qmax(self, x)
```

Return maximal action value for given state

$$\max_{u_i} \mathcal{Q}(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- `x`: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

## ValueFunction.Qmean

```
fitr.agents.value_functions.Qmean(self, x)
```

Return mean action value for given state

$$\text{Mean}(\mathcal{Q}(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|} \mathbf{1}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- `x`: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### ValueFunction.Qx

`fitr.agents.value_functions.Qx(self, x)`

Compute action values for a given state

$$Q(\mathbf{x}, :) = \mathbf{Q}\mathbf{x}$$

Arguments:

- `x`: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

### ValueFunction.Vx

`fitr.agents.value_functions.Vx(self, x)`

Compute value of state `x`

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- `x`: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state `x`

---

### ValueFunction.uQx

`fitr.agents.value_functions.uQx(self, u, x)`

Compute value of taking action `u` in state `x`

$$Q(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

### ValueFunction.update

```
fitr.agents.value_functions.update(self, x, u, r, x_, u_)
```

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector
  - **r**: Scalar reward
  - **x\_**: `ndarray((nstates,))` one-hot next-state vector
  - **u\_**: `ndarray((nactions,))` one-hot next-action vector
- 

### DummyLearner

```
fitr.agents.value_functions.DummyLearner()
```

A critic/value function for the random learner

This class actually contributes nothing except identifying that a value function has been chosen for an `Agent` object

Arguments:

- **env**: A `fitr.environments.Graph`
- 

### DummyLearner.Qmax

```
fitr.agents.value_functions.Qmax(self, x)
```

Return maximal action value for given state

$$\max_{u_i} Q(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^T \mathbf{Q} \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### DummyLearner.Qmean

```
fitr.agents.value_functions.Qmean(self, x)
```

Return mean action value for given state

$$\text{Mean}(Q(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|} \mathbf{1}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: ndarray((nstates,)) one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### DummyLearner.Qx

```
fitr.agents.value_functions.Qx(self, x)
```

Compute action values for a given state

$$Q(\mathbf{x}, :) = \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: ndarray((nstates,)) one-hot state vector

Returns:

ndarray((nactions,)) vector of values for actions in the given state

---

### DummyLearner.Vx

```
fitr.agents.value_functions.Vx(self, x)
```

Compute value of state **x**

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state **x**

---

### **DummyLerner.uQx**

```
fitr.agents.value_functions.uQx(self, u, x)
```

Compute value of taking action **u** in state **x**

$$Q(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

### **DummyLerner.update**

```
fitr.agents.value_functions.update(self, x, u, r, x_, u_)
```

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector
  - **r**: Scalar reward
  - **x\_**: `ndarray((nstates,))` one-hot next-state vector
  - **u\_**: `ndarray((nactions,))` one-hot next-action vector
- 

### **InstrumentalRescorlaWagnerLerner**

```
fitr.agents.value_functions.InstrumentalRescorlaWagnerLerner()
```

Learns an instrumental control policy through one-step error-driven updates of the state-action value function

The instrumental Rescorla-Wagner rule is as follows:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r - \mathbf{u}^\top \mathbf{Q}\mathbf{x})\mathbf{u}\mathbf{x}^\top,$$

where  $0 < \alpha < 1$  is the learning rate, and where the reward prediction error (RPE) is  $\delta = (r - \mathbf{u}^\top \mathbf{Q}\mathbf{x})$ .

\$\$

Arguments:

- **env**: A `fitr.environments.Graph`
  - **learning\_rate**: Learning rate  $\alpha$
- 

### **InstrumentalRescorlaWagnerLearner.Qmax**

`fitr.agents.value_functions.Qmax(self, x)`

Return maximal action value for given state

$$\max_{u_i} Q(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### **InstrumentalRescorlaWagnerLearner.Qmean**

`fitr.agents.value_functions.Qmean(self, x)`

Return mean action value for given state

$$Mean(Q(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|} \mathbf{1}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

**InstrumentalRescorlaWagnerLearner.Qx**

```
fitr.agents.value_functions.Qx(self, x)
```

Compute action values for a given state

$$Q(\mathbf{x}, :) = \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: ndarray((nstates,)) one-hot state vector

Returns:

ndarray((nactions,)) vector of values for actions in the given state

---

**InstrumentalRescorlaWagnerLearner.Vx**

```
fitr.agents.value_functions.Vx(self, x)
```

Compute value of state **x**

$$V(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- **x**: ndarray((nstates,)) one-hot state vector

Returns:

Scalar value of state **x**

---

**InstrumentalRescorlaWagnerLearner.uQx**

```
fitr.agents.value_functions.uQx(self, u, x)
```

Compute value of taking action **u** in state **x**

$$Q(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **u**: ndarray((nactions,)) one-hot action vector
- **x**: ndarray((nstates,)) one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

**InstrumentalRescorlaWagnerLearner.update**

```
fitr.agents.value_functions.update(self, x, u, r, x_, u_)
```

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector
  - **r**: Scalar reward
  - **x\_**: `ndarray((nstates,))` one-hot next-state vector
  - **u\_**: `ndarray((nactions,))` one-hot next-action vector
- 

**QLearner**

```
fitr.agents.value_functions.QLearner()
```

Learns an instrumental control policy through Q-learning

The Q-learning rule is as follows:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^T \mathbf{Q} \mathbf{x}' - \mathbf{u}^T \mathbf{Q} \mathbf{x}) \mathbf{z},$$

where  $0 < \alpha < 1$  is the learning rate,  $0 \leq \gamma \leq 1$  is a discount factor, and where the reward prediction error (RPE) is  $\delta = (r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^T \mathbf{Q} \mathbf{x}' - \mathbf{u}^T \mathbf{Q} \mathbf{x})$ . We have also included an eligibility trace  $\mathbf{z}$  defined as

$$\mathbf{z} = \mathbf{u} \mathbf{x}^T + \gamma \lambda \mathbf{z}$$

Arguments:

- **env**: A `fitr.environments.Graph`
  - **learning\_rate**: Learning rate  $\alpha$
  - **discount\_factor**: Discount factor  $\gamma$
  - **trace\_decay**: Eligibility trace decay  $\lambda$
- 

**QLearner.Qmax**

```
fitr.agents.value_functions.Qmax(self, x)
```

Return maximal action value for given state

$$\max_{u_i} Q(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^T \mathbf{Q} \mathbf{x}$$



Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### **QLearner.Qmean**

`fitr.agents.value_functions.Qmean(self, x)`

Return mean action value for given state

$$\text{Mean}(Q(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|} \mathbf{1}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### **QLearner.Qx**

`fitr.agents.value_functions.Qx(self, x)`

Compute action values for a given state

$$Q(\mathbf{x}, :) = \mathbf{Q} \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

### **QLearner.Vx**

`fitr.agents.value_functions.Vx(self, x)`

Compute value of state  $\mathbf{x}$

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- $\mathbf{x}$ : `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state  $\mathbf{x}$

---

### **QLearner.uQx**

```
fitr.agents.value_functions.uQx(self, u, x)
```

Compute value of taking action  $\mathbf{u}$  in state  $\mathbf{x}$

$$\mathcal{Q}(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- $\mathbf{u}$ : `ndarray((nactions,))` one-hot action vector
- $\mathbf{x}$ : `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action  $\mathbf{u}$  in state  $\mathbf{x}$

---

### **QLearner.update**

```
fitr.agents.value_functions.update(self, x, u, r, x_, u_)
```

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- $\mathbf{x}$ : `ndarray((nstates,))` one-hot state vector
  - $\mathbf{u}$ : `ndarray((nactions,))` one-hot action vector
  - $\mathbf{r}$ : Scalar reward
  - $\mathbf{x}_$ : `ndarray((nstates,))` one-hot next-state vector
  - $\mathbf{u}_$ : `ndarray((nactions,))` one-hot next-action vector
-

**SARSA Learner**

```
fitr.agents.value_functions.SARSA Learner()
```

Learns an instrumental control policy through the SARSA learning rule

The SARSA learning rule is as follows:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r + \gamma \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x}) \mathbf{z},$$

where  $0 < \alpha < 1$  is the learning rate,  $0 \leq \gamma \leq 1$  is a discount factor, and where the reward prediction error (RPE) is  $\delta = (r + \gamma \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$ . We have also included an eligibility trace  $\mathbf{z}$  defined as

$$\mathbf{z} = \mathbf{u} \mathbf{x}^\top + \gamma \lambda \mathbf{z}$$

Arguments:

- **env**: A `fitr.environments.Graph`
  - **learning\_rate**: Learning rate  $\alpha$
  - **discount\_factor**: Discount factor  $\gamma$
  - **trace\_decay**: Eligibility trace decay  $\lambda$
- 

**SARSA Learner.Qmax**

```
fitr.agents.value_functions.Qmax(self, x)
```

Return maximal action value for given state

$$\max_{u_i} Q(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

**SARSA Learner.Qmean**

```
fitr.agents.value_functions.Qmean(self, x)
```

Return mean action value for given state

$$Mean(Q(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|} \mathbf{1}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- `x`: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### **SARSA.Learner.Qx**

```
fitr.agents.value_functions.Qx(self, x)
```

Compute action values for a given state

$$Q(\mathbf{x}, :) = \mathbf{Q}\mathbf{x}$$

Arguments:

- `x`: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

### **SARSA.Learner.Vx**

```
fitr.agents.value_functions.Vx(self, x)
```

Compute value of state `x`

$$V(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- `x`: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state `x`

---

### **SARSA.Learner.uQx**

```
fitr.agents.value_functions.uQx(self, u, x)
```

Compute value of taking action `u` in state `x`

$$Q(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

### **SARSA Learner.update**

```
fitr.agents.value_functions.update(self, x, u, r, x_, u_)
```

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector
  - **r**: Scalar reward
  - **x\_**: `ndarray((nstates,))` one-hot next-state vector
  - **u\_**: `ndarray((nactions,))` one-hot next-action vector
- 

### **Agent**

```
fitr.agents.agents.Agent()
```

Base class for synthetic RL agents.

Arguments:

**meta** : List of metadata of arbitrary type. e.g. labels, covariates, etc. **params** : List of parameters for the agent. Should be filled for specific agent.

---

### **Agent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
-

**Agent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
  - **action**: `ndarray((nactions,))` one-hot action vector
  - **reward**: scalar reward
  - **next\_state**: `ndarray((nstates,))` one-hot next-state vector
  - **next\_action**: `ndarray((nactions,))` one-hot action vector
- 

**Agent.reset\_trace**

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector (optional)
- 

**BanditAgent**

```
fitr.agents.agents.BanditAgent()
```

A base class for agents in bandit tasks (i.e. with one step).

Arguments:

- **task**: `fitr.environments.Graph`
- 

**BanditAgent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
-

**BanditAgent.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials**: int number of trials

Returns:

```
fitr.data.BehaviouralData
```

---

**BanditAgent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: ndarray((nstates,)) one-hot state vector
  - **action**: ndarray((nactions,)) one-hot action vector
  - **reward**: scalar reward
  - **next\_state**: ndarray((nstates,)) one-hot next-state vector
  - **next\_action**: ndarray((nactions,)) one-hot action vector
- 

**BanditAgent.log\_prob**

```
fitr.agents.agents.log_prob(self, state)
```

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state**: ndarray((nstates,)) one-hot state vector

Returns:

```
ndarray((nactions,)) log-likelihood vector
```

---

**BanditAgent.reset\_trace**

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector (optional)
- 

**MDPAgent**

```
fitr.agents.agents.MDPAgent()
```

A base class for agents that operate on MDPs.

This mainly has implications for generating data.

Arguments:

- **task**: `fitr.environments.Graph`
- 

**MDPAgent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- 

**MDPAgent.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

```
fitr.data.BehaviouralData
```

---



**MDPAgent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: ndarray((nstates,)) one-hot state vector
  - **action**: ndarray((nactions,)) one-hot action vector
  - **reward**: scalar reward
  - **next\_state**: ndarray((nstates,)) one-hot next-state vector
  - **next\_action**: ndarray((nactions,)) one-hot action vector
- 

**MDPAgent.reset\_trace**

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: ndarray((nstates,)) one-hot state vector
  - **u**: ndarray((nactions,)) one-hot action vector (optional)
- 

**RandomBanditAgent**

```
fitr.agents.agents.RandomBanditAgent()
```

An agent that simply selects random actions at each trial

---

**RandomBanditAgent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: ndarray((nstates,)) one-hot state vector
-

**RandomBanditAgent.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials**: int number of trials

Returns:

```
fitr.data.BehaviouralData
```

---

**RandomBanditAgent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: ndarray((nstates,)) one-hot state vector
  - **action**: ndarray((nactions,)) one-hot action vector
  - **reward**: scalar reward
  - **next\_state**: ndarray((nstates,)) one-hot next-state vector
  - **next\_action**: ndarray((nactions,)) one-hot action vector
- 

**RandomBanditAgent.log\_prob**

```
fitr.agents.agents.log_prob(self, state)
```

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state**: ndarray((nstates,)) one-hot state vector

Returns:

```
ndarray((nactions,)) log-likelihood vector
```

---

**RandomBanditAgent.reset\_trace**

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector (optional)
- 

**RandomMDPAgent**

```
fitr.agents.agents.RandomMDPAgent()
```

An agent that simply selects random actions at each trial

**Notes**

This has been specified as an `OnPolicyAgent` arbitrarily.

---

**RandomMDPAgent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- 

**RandomMDPAgent.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

```
fitr.data.BehaviouralData
```

---

**RandomMDPAgent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state:** ndarray((nstates,)) one-hot state vector
  - **action:** ndarray((nactions,)) one-hot action vector
  - **reward:** scalar reward
  - **next\_state:** ndarray((nstates,)) one-hot next-state vector
  - **next\_action:** ndarray((nactions,)) one-hot action vector
- 

**RandomMDPAgent.reset\_trace**

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x:** ndarray((nstates,)) one-hot state vector
  - **u:** ndarray((nactions,)) one-hot action vector (optional)
- 

**SARSA SoftmaxAgent**

```
fitr.agents.agents.SARSA SoftmaxAgent()
```

An agent that uses the SARSA learning rule and a softmax policy

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i e^{\beta v_i}}.$$

The value function is SARSA:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r + \gamma \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x}) \mathbf{z},$$

where  $0 < \alpha < 1$  is the learning rate,  $0 \leq \gamma \leq 1$  is a discount factor, and where the reward prediction error (RPE) is  $\delta = (r + \gamma \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$ . We have also included an eligibility trace  $\mathbf{z}$  defined as

$$\mathbf{z} = \mathbf{u}\mathbf{x}^\top + \gamma\lambda\mathbf{z}$$

Arguments:

- **task**: `fitr.environments.Graph`
  - **learning\_rate**: Learning rate  $\alpha$
  - **discount\_factor**: Discount factor  $\gamma$
  - **trace\_decay**: Eligibility trace decay  $\lambda$
  - **inverse\_softmax\_temp**: Inverse softmax temperature  $\beta$
  - **rng**: `np.random.RandomState`
- 

### **SARSA SoftmaxAgent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- 

### **SARSA SoftmaxAgent.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

```
fitr.data.BehaviouralData
```

---

### **SARSA SoftmaxAgent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

- **action:** `ndarray((nactions,))` one-hot action vector
  - **reward:** scalar reward
  - **next\_state:** `ndarray((nstates,))` one-hot next-state vector
  - **next\_action:** `ndarray((nactions,))` one-hot action vector
- 

### SARSA Softmax Agent `reset_trace`

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x:** `ndarray((nstates,))` one-hot state vector
  - **u:** `ndarray((nactions,))` one-hot action vector (optional)
- 

### QLearning Softmax Agent

```
fitr.agents.agents.QLearningSoftmaxAgent()
```

An agent that uses the Q-learning rule and a softmax policy

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i e^{\beta v_i}}.$$

The value function is Q-learning:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha (r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x}) \mathbf{z},$$

where  $0 < \alpha < 1$  is the learning rate,  $0 \leq \gamma \leq 1$  is a discount factor, and where the reward prediction error (RPE) is  $\delta = (r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$ . The eligibility trace  $\mathbf{z}$  is defined as

$$\mathbf{z} = \mathbf{u} \mathbf{x}^\top + \gamma \lambda \mathbf{z}$$

Arguments:

- **task:** `fitr.environments.Graph`
- **learning\_rate:** Learning rate  $\alpha$
- **discount\_factor:** Discount factor  $\gamma$

- **trace\_decay**: Eligibility trace decay  $\lambda$
  - **inverse\_softmax\_temp**: Inverse softmax temperature  $\beta$
  - **rng**: `np.random.RandomState`
- 

### **QLearningSoftmaxAgent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- 

### **QLearningSoftmaxAgent.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

```
fitr.data.BehaviouralData
```

---

### **QLearningSoftmaxAgent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
  - **action**: `ndarray((nactions,))` one-hot action vector
  - **reward**: scalar reward
  - **next\_state**: `ndarray((nstates,))` one-hot next-state vector
  - **next\_action**: `ndarray((nactions,))` one-hot action vector
-

**QLearningSoftmaxAgent.reset\_trace**

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: ndarray (nstates, ) one-hot state vector
  - **u**: ndarray (nactions, ) one-hot action vector (optional)
- 

**RWSoftmaxAgent**

```
fitr.agents.agents.RWSoftmaxAgent()
```

An instrumental Rescorla-Wagner agent with a softmax policy

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i e^{\beta v_i}}.$$

The value function is the Rescorla-Wagner learning rule:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r - \mathbf{u}^\top \mathbf{Q} \mathbf{x}) \mathbf{u} \mathbf{x}^\top,$$

where  $0 < \alpha < 1$  is the learning rate,  $0 \leq \gamma \leq 1$  is a discount factor, and where the reward prediction error (RPE) is  $\delta = (r - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$ .

Arguments:

- **task**: fitr.environments.Graph
  - **learning\_rate**: Learning rate  $\alpha$
  - **inverse\_softmax\_temp**: Inverse softmax temperature  $\beta$
  - **rng**: np.random.RandomState
- 

**RWSoftmaxAgent.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:



- **state:** `ndarray((nstates,))` one-hot state vector
- 

### **RWSoftmaxAgent.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials:** `int` number of trials

Returns:

```
fitr.data.BehaviouralData
```

---

### **RWSoftmaxAgent.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state:** `ndarray((nstates,))` one-hot state vector
  - **action:** `ndarray((nactions,))` one-hot action vector
  - **reward:** scalar reward
  - **next\_state:** `ndarray((nstates,))` one-hot next-state vector
  - **next\_action:** `ndarray((nactions,))` one-hot action vector
- 

### **RWSoftmaxAgent.log\_prob**

```
fitr.agents.agents.log_prob(self, state)
```

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state:** `ndarray((nstates,))` one-hot state vector

Returns:

```
ndarray((nactions,)) log-likelihood vector
```

---

**RWSoftmaxAgent.reset\_trace**

```
fitr.agents.agents.reset_trace(self, x, u=None)
```

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: ndarray (nstates,) one-hot state vector
  - **u**: ndarray (nactions,) one-hot action vector (optional)
- 

**RWSoftmaxAgentRewardSensitivity**

```
fitr.agents.agents.RWSoftmaxAgentRewardSensitivity()
```

An instrumental Rescorla-Wagner agent with a softmax policy, whose experienced reward is scaled by a factor  $\rho$ .

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i e^{\beta v_i}}.$$

The value function is the Rescorla-Wagner learning rule with scaled reward  $\rho r$ :

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(\rho r - \mathbf{u}^\top \mathbf{Q} \mathbf{x}) \mathbf{u} \mathbf{x}^\top,$$

where  $0 < \alpha < 1$  is the learning rate,  $0 \leq \gamma \leq 1$  is a discount factor, and where the reward prediction error (RPE) is  $\delta = (\rho r - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$ .

Arguments:

- **task**: fitr.environments.Graph
  - **learning\_rate**: Learning rate  $\alpha$
  - **inverse\_softmax\_temp**: Inverse softmax temperature  $\beta$
  - **reward\_sensitivity**: Reward sensitivity parameter  $\rho$
  - **rng**: np.random.RandomState
- 

**RWSoftmaxAgentRewardSensitivity.action**

```
fitr.agents.agents.action(self, state)
```

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state:** `ndarray((nstates,))` one-hot state vector
- 

### **RWSoftmaxAgentRewardSensitivity.generate\_data**

```
fitr.agents.agents.generate_data(self, ntrials)
```

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials:** `int` number of trials

Returns:

```
fitr.data.BehaviouralData
```

---

### **RWSoftmaxAgentRewardSensitivity.learning**

```
fitr.agents.agents.learning(self, state, action, reward, next_state, next_action)
```

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state:** `ndarray((nstates,))` one-hot state vector
  - **action:** `ndarray((nactions,))` one-hot action vector
  - **reward:** scalar reward
  - **next\_state:** `ndarray((nstates,))` one-hot next-state vector
  - **next\_action:** `ndarray((nactions,))` one-hot action vector
- 

### **RWSoftmaxAgentRewardSensitivity.log\_prob**

```
fitr.agents.agents.log_prob(self, state)
```

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state:** `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` log-likelihood vector

---

### **RWSoftmaxAgentRewardSensitivity.reset\_trace**

`fitr.agents.agents.reset_trace(self, x, u=None)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
  - **u**: `ndarray((nactions,))` one-hot action vector (optional)
-

# Chapter 5

## Data

### `fitr.data`

A module containing a generic class for behavioural data.

#### **BehaviouralData**

```
fitr.data.BehaviouralData()
```

A flexible and generic object to store and process behavioural data across tasks

Arguments:

- **ngroups**: Integer number of groups represented in the dataset. Only > 1 if data are merged
  - **nsubjects**: Integer number of subjects in dataset
  - **ntrials**: Integer number of trials done by each subject
  - **dict**: Dictionary storage indexed by subject.
  - **params**: `ndarray(nsubjects, nparams + 1)` parameters for each (simulated) subject
  - **meta**: Array of covariates of type `ndarray(nsubjects, nmetadata_features+1)`
  - **tensor**: Tensor representation of the behavioural data of type `ndarray(nsubjects, ntrials, nfeatures)`
- 

#### **BehaviouralData.add\_subject**

```
fitr.data.add_subject(self, subject_index, parameters, subject_meta)
```

Appends a new subject to the dataset

Arguments:

- **subject\_index**: Integer identification for subject
  - **parameters**: `list` of parameters for the subject
  - **subject\_meta**: Some covariates for the subject (`list`)
-

**BehaviouralData.initialize\_data\_dictionary**

```
fitr.data.initialize_data_dictionary(self)
```

---

**BehaviouralData.make\_behavioural\_ngrams**

```
fitr.data.make_behavioural_ngrams(self, n)
```

Creates N-grams of behavioural data

---

**BehaviouralData.make\_cooccurrence\_matrix**

```
fitr.data.make_cooccurrence_matrix(self, k, dtype=<class 'numpy.float32'>)
```

---

**BehaviouralData.make\_tensor\_representations**

```
fitr.data.make_tensor_representations(self)
```

Creates a tensor with all subjects' data

**Notes**

Assumes that all subjects did same number of trials.

---

**BehaviouralData.numpy\_tensor\_to\_bdf**

```
fitr.data.numpy_tensor_to_bdf(self, X)
```

Creates BehaviouralData formatted set from a dataset stored in a numpy ndarray.

Arguments:

- **X**: ndarray((nsubjects, ntrials, m)) with m being the size of flattened single-trial data
- 

**BehaviouralData.unpack\_tensor**

```
fitr.data.unpack_tensor(self, x_dim, u_dim, r_dim=1, terminal_dim=1, get='sarsat')
```

Unpacks data stored in tensor format into separate arrays for states, actions, rewards, next states, and next actions.

Arguments:

`x_dim` : Task state space dimensionality (`int`) `u_dim` : Task action space dimensionality (`int`) `r_dim` : Reward dimensionality (`int`, default=1) `terminal_dim` : Dimensionality of the terminal state indicator (`int`, default=1) `get` : String indicating the order that data are stored in the array. Can also be shortened such that fewer elements are returned. For example, the default is `sarsat`.

Returns:

List with data, where each element is in the order of the argument `get`

---

### **BehaviouralData.update**

```
fitr.data.update(self, subject_index, behav_data)
```

Adds behavioural data to the dataset

Arguments:

- **subject\_index**: Integer index for the subject
  - **behav\_data**: 1-dimensional ndarray of flattened data
- 

### **merge\_behavioural\_data**

```
fitr.data.merge_behavioural_data(datalist)
```

Combines BehaviouralData objects.

Arguments:

- **datalist**: List of BehaviouralData objects

Returns:

BehaviouralData with data from multiple groups merged.

---

# Chapter 6

## Inference

### `fitr.inference`

Methods for inferring the parameters of generative models for reinforcement learning data.

#### **OptimizationResult**

`fitr.inference.optimization_result.OptimizationResult()`

Container for the results of an optimization run on a generative model of behavioural data

Arguments:

- **subject\_id**: `ndarray((nsubjects,))` or `None` (default). Integer ids for subjects
  - **xmin**: `ndarray((nsubjects,nparams))` or `None` (default). Parameters that minimize objective function
  - **fmin**: `ndarray((nsubjects,))` or `None` (default). Value of objective function at minimum
  - **fevals**: `ndarray((nsubjects,))` or `None` (default). Number of function evaluations required to minimize objective function
  - **niters**: `ndarray((nsubjects,))` or `None` (default). Number of iterations required to minimize objective function
  - **lme**: `ndarray((nsubjects,))` or `None` (default). Log model evidence
  - **bic**: `ndarray((nsubjects,))` or `None` (default). Bayesian Information Criterion
  - **hess\_inv**: `ndarray((nsubjects,nparams,nparams))` or `None` (default). Inverse Hessian at the optimum.
  - **err**: `ndarray((nsubjects,nparams))` or `None` (default). Error of estimates at optimum.
- 

#### **OptimizationResult.transform\_xmin**

`fitr.inference.optimization_result.transform_xmin(self, transforms, inplace=False)`

Rescales the parameter estimates.

Arguments:



- **transforms**: list. Transformation functions where `len(transforms) == self.xmin.shape[1]`
- **inplace**: bool. Whether to change the values in `self.xmin`. Default is `False`, which returns an `ndarray((nsubjects, nparams))` of the transformed parameters.

Returns:

`ndarray((nsubjects, nparams))` of the transformed parameters if `inplace=False`

---

## mlepar

```
fitr.inference.mle_parallel.mlepar(f, data, nparams, minstarts=2, maxstarts=10, init
```

Computes maximum likelihood estimates using parallel CPU resources.

Wraps over the `fitr.optimization.mle_parallel.mle` function.

Arguments:

- **f**: Likelihood function
- **data**: A subscriptable object whose first dimension indexes subjects
- **optimizer**: Optimization function (currently only `l_bfgs_b` supported)
- **nparams**: int number of parameters to be estimated
- **minstarts**: int. Minimum number of restarts with new initial values
- **maxstarts**: int. Maximum number of restarts with new initial values
- **init\_sd**: Standard deviation for Gaussian initial values

Returns:

`fitr.inference.OptimizationResult`

---

## l\_bfgs\_b

```
fitr.inference.mle_parallel.l_bfgs_b(f, i, data, nparams, minstarts=2, maxstarts=10,
```

Minimizes the negative log-probability of data with respect to some parameters under function `f` using the L-BFGS-B algorithm.

This function is specified for use with parallel CPU resources.

Arguments:

- **f**: Log likelihood function
- **i**: int. Subject being optimized (slices first dimension of data)
- **data**: Object subscriptable along first dimension to indicate subject being optimized
- **nparams**: int. Number of parameters in the model
- **minstarts**: int. Minimum number of restarts with new initial values
- **maxstarts**: int. Maximum number of restarts with new initial values
- **init\_sd**: Standard deviation for Gaussian initial values

Returns:

- **i**: `int`. Subject being optimized (slices first dimension of data)
  - **xmin**: `ndarray((nparams,))`. Parameter values at optimum
  - **fmin**: Scalar objective function value at optimum
  - **fevals**: `int`. Number of function evaluations
  - **niters**: `int`. Number of iterations
  - **lme\_**: Scalar log-model evidence at optimum
  - **bic\_**: Scalar Bayesian Information Criterion at optimum
  - **hess\_inv**: `ndarray((nparams, nparams))`. Inv at optimum
-

## Chapter 7

# Criticism

### **fitr.criticism**

Methods for criticism of model fits.

### **actual\_estimate**

```
fitr.criticism.plotting.actual_estimate(y_true, y_pred, xlabel='Actual', ylabel='Est
```

Plots parameter estimates against the ground truth values.

Arguments:

- **y\_true**: ndarray(nsamples). Vector of ground truth parameters
- **y\_pred**: ndarray(nsamples). Vector of parameter estimates
- **xlabel**: str. Label for x-axis
- **ylabel**: str. Label for y-axis
- **corr**: bool. Whether to plot correlation coefficient.
- **figsize**: tuple. Figure size (inches).

Returns:

```
matplotlib.pyplot.Figure
```

---

## Chapter 8

# Metrics

### **fitr.metrics**

Metrics and performance statistics.

#### **bic**

```
fitr.metrics.bic(log_prob, nparams, ntrials)
```

Bayesian Information Criterion (BIC)

Arguments:

- **log\_prob**: Log probability
- **nparams**: Number of parameters in the model
- **ntrials**: Number of trials in the time series

Returns:

Scalar estimate of BIC.

---

#### **linear\_correlation**

```
fitr.metrics.linear_correlation(X, Y)
```

Linear correlation coefficient.

Will compute the following formula

$$\rho = \frac{\mathbf{x}^\top \mathbf{y}}{\|\mathbf{x}_{Vert}\| \cdot \|\mathbf{y}_{Vert}\|}$$

where each vector  $\mathbf{x}$  and  $\mathbf{y}$  are rows of the matrices  $\mathbf{X}$  and  $\mathbf{Y}$ , respectively.

Arguments:

- **X**: `ndarray((nsamples, nfeatures))` of dimension 1 or 2. If X is a 1D array, it will be converted to 2D prior to computation
- **Y**: `ndarray((nsamples, nfeatures))` of dimension 1 or 2. If Y is a 1D array, it will be converted to 2D prior to computation

Returns:

- **rho**: `ndarray((nfeatures,))`. Correlation coefficient(s)

TODO:

- [ ] Create error raised when X and Y are not same dimension
- 

## **lme**

`fitr.metrics.lme(log_prob, nparams, hess_inv)`

Laplace approximation to the log model evidence

Arguments:

- **log\_prob**: Log probability
- **nparams**: Number of parameters in the model
- **hess\_inv**: Hessian at the optimum (shape is  $K \times K$ )

Returns:

Scalar approximation of the log model evidence

---

## **log\_loss**

`fitr.metrics.log_loss(p, q)`

Computes log loss.

$$\mathcal{L} = \mathbf{p}^\top \log \mathbf{q} + (1 - \mathbf{p})^\top \log(1 - \mathbf{q})$$

Arguments:

- **p**: Binary vector of true labels `ndarray((nsamples,))`
- **q**: Vector of estimates (between 0 and 1) of type `ndarray((nsamples,))`

Returns:

Scalar log loss

---

# Chapter 9

## Utilities

### `fitr.utils`

Functions used across `fitr`.

#### `logsumexp`

`fitr.utils.logsumexp(x)`

Numerically stable `logsumexp`.

Computed as follows:

$$\max x + \log \sum_x e^{x - \max x}$$

Arguments:

- `x`: ‘ndarray(shape=(nactions,))’

Returns:

float

---

#### `relu`

`fitr.utils.relu(x, a_max=None)`

Rectified linearity

$$\mathbf{x}' = \max(x_i, 0)_{i=1}^{|\mathbf{x}|}$$

Arguments:

- `x`: Vector of inputs

- **a\_max**: Upper bound at which to clip values of  $\mathbf{x}$

Returns:

Exponentiated values of  $\mathbf{x}$ .

---

## scale\_data

```
fitr.utils.scale_data(X, axis=0, with_mean=True, with_var=True)
```

Rescales data by subtracting mean and dividing by variance

$$\mathbf{x}' = \frac{\mathbf{x} - \frac{1}{n} \mathbf{1}^\top \mathbf{x}}{\text{Var}(\mathbf{x})}$$

Arguments:

- **X**: `ndarray((nsamples, [nfeatures]))`. Data. May be 1D or 2D.
- **with\_mean**: `bool`. Whether to subtract the mean
- **with\_var**: `bool`. Whether to divide by variance

Returns:

`ndarray(X.shape)`. Rescaled data.

---

## sigmoid

```
fitr.utils.sigmoid(x, a_min=-10, a_max=10)
```

Sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Arguments:

- **x**: Vector
- **a\_min**: Lower bound at which to clip values of  $\mathbf{x}$
- **a\_max**: Upper bound at which to clip values of  $\mathbf{x}$

Returns:

Vector between 0 and 1 of size `x.shape`

---

**softmax**

```
fitr.utils.softmax(x)
```

Computes the softmax function

$$p(\mathbf{x}) = \frac{e^{\mathbf{x} - \max_i x_i}}{\mathbf{1}^\top e^{\mathbf{x} - \max_i x_i}}$$

Arguments:

- **x**: Softmax logits (`ndarray (N, )`)

Returns:

Vector of probabilities of size `ndarray (N, )`

---

**stable\_exp**

```
fitr.utils.stable_exp(x, a_min=-10, a_max=10)
```

Clipped exponential function

Avoids overflow by clipping input values.

Arguments:

- **x**: Vector of inputs
- **a\_min**: Lower bound at which to clip values of `x`
- **a\_max**: Upper bound at which to clip values of `x`

Returns:

Exponentiated values of `x`.

---