# fitr

Abraham Nunes

July 2, 2018

# Contents

# Chapter 1

# Overview & Foundations

# Chapter 2

# Tutorials

## Getting Started

### Installation

```
pip install git+https://github.com/abrahamnunes/fitr.git
```

## Simulating and Fitting a Two-Armed Bandit

```python
import numpy as np
import matplotlib.pyplot as plt
from fitr import generate_behavioural_data
from fitr.environments import TwoArmedBandit
from fitr.agents import RWSoftmaxAgent
from fitr.inference import mlepar
from fitr.utils import sigmoid
from fitr.utils import relu
from fitr.criticism.plotting import actual_estimate


N = 50  # number of subjects
T = 200 # number of trials

# Generate synthetic data
data = generate_behavioural_data(TwoArmedBandit, RWSoftmaxAgent, N, T)

# Create log-likelihood function
def log_prob(w, D):
    lr  = sigmoid(w[0], a_min=-6, a_max=6)
    ist = relu(w[1], a_max=10)
    agent = RWSoftmaxAgent(TwoArmedBandit(), lr, ist)
    L = 0
    for t in range(D.shape[0]):
        x=D[t,:3]; u=D[t,3:5]; r=D[t,5]; x_=D[t,6:]
```

```python
        L += u@agent.log_prob(x)
        agent.learning(x, u, r, x_, None)
    return L

# Fit model
res = mlepar(log_prob, data.tensor, nparams=2, maxstarts=5)
X = res.transform_xmin([sigmoid, relu])

# Criticism: Actual vs. Estimate Plots
lr_fig  = actual_estimate(data.params[:,1], X[:,0]); plt.show()
ist_fig = actual_estimate(data.params[:,2], X[:,1]); plt.show()
```

## Simulating and Fitting Data from a Random Contextual Bandit Task

```python
import numpy as np
import matplotlib.pyplot as plt
from fitr import generate_behavioural_data
from fitr.agents import RWSoftmaxAgent
from fitr.environments import RandomContextualBandit
from fitr.criticism.plotting import actual_estimate
from fitr.inference import mlepar
from fitr.utils import sigmoid, relu

class MyBanditTask(RandomContextualBandit):
    def __init__(self):
        super().__init__(nactions=4,
                         noutcomes=3,
                         nstates=4,
                         min_actions_per_context=None,
                         alpha=0.1,
                         alpha_start=1.,
                         shift_flip='shift',
                         reward_lb=-1,
                         reward_ub=1,
                         reward_drift='on',
                         drift_mu=np.zeros(3),
                         drift_sd=1.)

data = generate_behavioural_data(MyBanditTask, RWSoftmaxAgent, 20, 200)

def log_prob(w, D):
    agent = RWSoftmaxAgent(task=MyBanditTask(),
                           learning_rate=w[0],
                           inverse_softmax_temp=w[1])
    L=0
    for t in range(D.shape[0]):
        x=D[t,:7]; u=D[t,7:11]; r=D[t,11]; x_=D[t,12:]
```

```
        L += u@agent.log_prob(x)
        agent.learning(x, u, r, x_, None)
    return L


res = mlepar(log_prob, data.tensor, 2, maxstarts=5)
X = res.transform_xmin([sigmoid, relu])

# Criticism: Actual vs. Estimate Plots
lr_fig  = actual_estimate(data.params[:,1], X[:,0]); plt.show()
ist_fig = actual_estimate(data.params[:,2], X[:,1]); plt.show()
```

# Part I

# API

# Chapter 3

# Environments

## **fitr.environments**

Functions to synthesize data from behavioural tasks.

### **Graph**

`fitr.environments.Graph()`

Base object that defines a reinforcement learning task.

#### **Definitions**

- $\mathbf{x} \in \mathcal{X}$ be a one-hot state vector, where $|\mathcal{X}| = n_x$
- $\mathbf{u} \in \mathcal{U}$ be a one-hot action vector, where $|\mathcal{U}| = n_u$
- $\mathsf{T} = p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$ be a transition tensor
- $p(\mathbf{x})$ be a distribution over starting states
- $\mathcal{J} : \mathcal{X} \to \mathcal{R}$, where $\mathcal{R} \subseteq \mathbb{R}$ be a reward function

Arguments:

- **T**: Transition tensor
- **R**: Vector of rewards for each state such that scalar reward $r_t = \mathbf{r}^o p\mathbf{x}$
- **end_states**: A vector $\{0, 1\}^{n_x}$ identifying which states terminate a trial (aka episode)
- **p_start**: Initial state distribution
- **label**: A string identifying a name for the task
- **state_labels**: A list or array of strings labeling the different states (for plotting purposes)
- **action_labels**: A list or array of strings labeling the different actions (for plotting purposes)
- **rng**: `np.random.RandomState` object
- **f_reward**: A function whose first argument is a vector of rewards for each state, and whose second argument is a state vector, and whose output is a scalar reward
- **cmap**: Matplotlib colormap for plotting.

**Notes**

There are two critical methods for the `Graph` class: `observation()` and `step`. All instances of a `Graph` must be able to call these functions. Let's say you have some bandit task `MyBanditTask` that inherits from `Graph`. To run such a task would look something like this:

```
env = MyBanditTask()            # Instantiate your environment object
agent = MyAgent()               # Some agent object (arbitrary, really)
for t in range(ntrials):
    x = env.observation()       # Samples initial state
    u = agent.action(x)         # Choose some action
    x_, r, done = agent.step(u) # Transition based on action
```

What differentiates tasks are the transition tensor $\mathsf{T}$, starting state distribution $p(\mathbf{x})$ and reward function $\mathcal{J}$ (which here would include the reward vector $\mathbf{r}$).

---

**Graph.adjacency_matrix_decomposition**

`fitr.environments.adjacency_matrix_decomposition(self)`

Singular value decomposition of the graph adjacency matrix

---

**Graph.get_graph_depth**

`fitr.environments.get_graph_depth(self)`

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**Graph.laplacian_matrix_decomposition**

`fitr.environments.laplacian_matrix_decomposition(self)`

Singular value decomposition of the graph Laplacian

---

**Graph.make_action_labels**

`fitr.environments.make_action_labels(self)`

Creates labels for the actions (for plotting) if none provided

---

### Graph.make_digraph

`fitr.environments.make_digraph(`self`)`

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

### Graph.make_state_labels

`fitr.environments.make_state_labels(`self`)`

Creates labels for the states (for plotting) if none provided

---

### Graph.make_undirected_graph

`fitr.environments.make_undirected_graph(`self`)`

Converts the DiGraph to undirected and computes some stats

---

### Graph.observation

`fitr.environments.observation(`self`)`

Samples an initial state from the start-state distribution $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

`x = env.observation()`

---

**Graph.plot_action_outcome_probabilities**

`fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None`

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

**Graph.plot_graph**

`fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.`

Plots the directed graph of the task

---

**Graph.plot_spectral_properties**

`fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile`

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

**Graph.random_action**

`fitr.environments.random_action(self)`

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\Big(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\Big)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

`u = env.random_action()`

---

**Graph.step**

`fitr.environments.step(`<span style="color:blue">`self`</span>`, action)`

Executes a state transition in the environment.

Arguments:

action : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

## TwoArmedBandit

`fitr.environments.TwoArmedBandit()`

A simple 2-armed bandit task

---

**TwoArmedBandit.adjacency_matrix_decomposition**

`fitr.environments.adjacency_matrix_decomposition(`<span style="color:blue">`self`</span>`)`

Singular value decomposition of the graph adjacency matrix

---

**TwoArmedBandit.get_graph_depth**

`fitr.environments.get_graph_depth(`<span style="color:blue">`self`</span>`)`

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

## TwoArmedBandit.laplacian_matrix_decomposition

`fitr.environments.laplacian_matrix_decomposition(`self`)`

Singular value decomposition of the graph Laplacian

---

## TwoArmedBandit.make_action_labels

`fitr.environments.make_action_labels(`self`)`

Creates labels for the actions (for plotting) if none provided

---

## TwoArmedBandit.make_digraph

`fitr.environments.make_digraph(`self`)`

Creates a `networkx` `DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

## TwoArmedBandit.make_state_labels

`fitr.environments.make_state_labels(`self`)`

Creates labels for the states (for plotting) if none provided

---

## TwoArmedBandit.make_undirected_graph

`fitr.environments.make_undirected_graph(`self`)`

Converts the DiGraph to undirected and computes some stats

---

## TwoArmedBandit.observation

`fitr.environments.observation(`self`)`

Samples an initial state from the start-state distribution $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

**TwoArmedBandit.plot_action_outcome_probabilities**

`fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None`

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

**TwoArmedBandit.plot_graph**

`fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.`

Plots the directed graph of the task

---

**TwoArmedBandit.plot_spectral_properties**

`fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile`

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

**TwoArmedBandit.random_action**

`fitr.environments.random_action(self)`

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\Big(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\Big)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

**TwoArmedBandit.step**

`fitr.environments.step(self, action)`

Executes a state transition in the environment.

Arguments:

action : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

**OrthogonalGoNoGo**

`fitr.environments.OrthogonalGoNoGo()`

The Orthogonal GoNogo task from Guitart-Masip et al. (2012)

---

**OrthogonalGoNoGo.adjacency_matrix_decomposition**

`fitr.environments.adjacency_matrix_decomposition(self)`

Singular value decomposition of the graph adjacency matrix

---

**OrthogonalGoNoGo.get_graph_depth**

`fitr.environments.get_graph_depth(self)`

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**OrthogonalGoNoGo.laplacian_matrix_decomposition**

`fitr.environments.laplacian_matrix_decomposition(`self`)`

Singular value decomposition of the graph Laplacian

---

**OrthogonalGoNoGo.make_action_labels**

`fitr.environments.make_action_labels(`self`)`

Creates labels for the actions (for plotting) if none provided

---

**OrthogonalGoNoGo.make_digraph**

`fitr.environments.make_digraph(`self`)`

Creates a `networkx` `DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

**OrthogonalGoNoGo.make_state_labels**

`fitr.environments.make_state_labels(`self`)`

Creates labels for the states (for plotting) if none provided

---

**OrthogonalGoNoGo.make_undirected_graph**

`fitr.environments.make_undirected_graph(`self`)`

Converts the DiGraph to undirected and computes some stats

---

**OrthogonalGoNoGo.observation**

`fitr.environments.observation(`self`)`

Samples an initial state from the start-state distribution $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

**OrthogonalGoNoGo.plot_action_outcome_probabilities**

`fitr.environments.plot_action_outcome_probabilities(`self`, figsize=`None`, outfile=`None`

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

**OrthogonalGoNoGo.plot_graph**

`fitr.environments.plot_graph(`self`, figsize=`None`, node_size=`2000`, arrowsize=`20`, lw=1.`

Plots the directed graph of the task

---

**OrthogonalGoNoGo.plot_spectral_properties**

`fitr.environments.plot_spectral_properties(`self`, figsize=`None`, outfile=`None`, outfile`

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

**OrthogonalGoNoGo.random_action**

`fitr.environments.random_action(`self`)`

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\Big(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\Big)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

### OrthogonalGoNoGo.step

`fitr.environments.step(`<span style="color:blue">`self`</span>`, action)`

Executes a state transition in the environment.

Arguments:

action : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

### TwoStep

`fitr.environments.TwoStep()`

An implementation of the Two-Step Task from Daw et al. (2011).

Arguments:

- **mu**: `float` identifying the drift of the reward-determining Gaussian random walks
- **sd**: `float` identifying the standard deviation of the reward-determining Gaussian random walks

---

### TwoStep.adjacency_matrix_decomposition

`fitr.environments.adjacency_matrix_decomposition(`<span style="color:blue">`self`</span>`)`

Singular value decomposition of the graph adjacency matrix

---

### TwoStep.f_reward

`fitr.environments.f_reward(`<span style="color:blue">`self`</span>`, R, x)`

---

**TwoStep.get_graph_depth**

`fitr.environments.get_graph_depth(`<span style="color:blue">`self`</span>`)`

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

**TwoStep.laplacian_matrix_decomposition**

`fitr.environments.laplacian_matrix_decomposition(`<span style="color:blue">`self`</span>`)`

Singular value decomposition of the graph Laplacian

---

**TwoStep.make_action_labels**

`fitr.environments.make_action_labels(`<span style="color:blue">`self`</span>`)`

Creates labels for the actions (for plotting) if none provided

---

**TwoStep.make_digraph**

`fitr.environments.make_digraph(`<span style="color:blue">`self`</span>`)`

Creates a `networkx DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

**TwoStep.make_state_labels**

`fitr.environments.make_state_labels(`<span style="color:blue">`self`</span>`)`

Creates labels for the states (for plotting) if none provided

---

**TwoStep.make_undirected_graph**

`fitr.environments.make_undirected_graph(self)`

Converts the DiGraph to undirected and computes some stats

---

**TwoStep.observation**

`fitr.environments.observation(self)`

Samples an initial state from the start-state distribution $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

`x = env.observation()`

---

**TwoStep.plot_action_outcome_probabilities**

`fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None`

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

**TwoStep.plot_graph**

`fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.`

Plots the directed graph of the task

---

**TwoStep.plot_reward_paths**

`fitr.environments.plot_reward_paths(self, outfile=None, outfiletype='pdf', figsize=N`

---

**TwoStep.plot_spectral_properties**

`fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile`

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

**TwoStep.random_action**

`fitr.environments.random_action(self)`

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\Big(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\Big)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

`u = env.random_action()`

---

**TwoStep.step**

`fitr.environments.step(self, action)`

Executes a state transition in the environment.

Arguments:

action : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

## ReverseTwoStep

`fitr.environments.ReverseTwoStep()`

From Kool & Gershman 2016.

---

### ReverseTwoStep.adjacency_matrix_decomposition

`fitr.environments.adjacency_matrix_decomposition(`self`)`

Singular value decomposition of the graph adjacency matrix

---

### ReverseTwoStep.f_reward

`fitr.environments.f_reward(`self`, R, x)`

---

### ReverseTwoStep.get_graph_depth

`fitr.environments.get_graph_depth(`self`)`

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

### ReverseTwoStep.laplacian_matrix_decomposition

`fitr.environments.laplacian_matrix_decomposition(`self`)`

Singular value decomposition of the graph Laplacian

---

### ReverseTwoStep.make_action_labels

`fitr.environments.make_action_labels(`self`)`

Creates labels for the actions (for plotting) if none provided

---

### ReverseTwoStep.make_digraph

`fitr.environments.make_digraph(self)`

Creates a `networkx` `DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

### ReverseTwoStep.make_state_labels

`fitr.environments.make_state_labels(self)`

Creates labels for the states (for plotting) if none provided

---

### ReverseTwoStep.make_undirected_graph

`fitr.environments.make_undirected_graph(self)`

Converts the DiGraph to undirected and computes some stats

---

### ReverseTwoStep.observation

`fitr.environments.observation(self)`

Samples an initial state from the start-state distribution $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

`x = env.observation()`

---

### ReverseTwoStep.plot_action_outcome_probabilities

`fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None`

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

### ReverseTwoStep.plot_graph

`fitr.environments.plot_graph(`self`, figsize=`None`, node_size=`2000`, arrowsize=`20`, lw=1.`

Plots the directed graph of the task

---

### ReverseTwoStep.plot_spectral_properties

`fitr.environments.plot_spectral_properties(`self`, figsize=`None`, outfile=`None`, outfile`

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

### ReverseTwoStep.random_action

`fitr.environments.random_action(`self`)`

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\Big(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\Big)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

`u = env.random_action()`

---

### ReverseTwoStep.step

`fitr.environments.step(`self`, action)`

Executes a state transition in the environment.

Arguments:

action : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

## RandomContextualBandit

`fitr.environments.RandomContextualBandit()`

Generates a random bandit task

Arguments:

- **nactions**: Number of actions
- **noutcomes**: Number of outcomes
- **nstates**: Number of contexts
- **min_actions_per_context**: Different contexts may have more or fewer actions than others (never more than `nactions`). This variable describes the minimum number of actions allowed in a context.
- **alpha**:
- **alpha_start**:
- **shift_flip**:
- **reward_lb**: Lower bound for drifting rewards
- **reward_ub**: Upper bound for drifting rewards
- **reward_drift**: Values (`on` or `off`) determining whether rewards are allowed to drift
- **drift_mu**: Mean of the Gaussian random walk determining reward
- **drift_sd**: Standard deviation of Gaussian random walk determining reward

---

## RandomContextualBandit.adjacency_matrix_decomposition

`fitr.environments.adjacency_matrix_decomposition(self)`

Singular value decomposition of the graph adjacency matrix

---

## RandomContextualBandit.f_reward

`fitr.environments.f_reward(self, R, x)`

---

## RandomContextualBandit.get_graph_depth

`fitr.environments.get_graph_depth(self)`

Returns the depth of the task graph.

Calculated as the depth from `START` (pre-initial state) to `END` (which absorbs trial from all terminal states), minus 2 to account for the `START->node` & `node->END` transitions.

Returns:

An `int` identifying the depth of the current graph for a single trial of the task

---

### RandomContextualBandit.laplacian_matrix_decomposition

`fitr.environments.laplacian_matrix_decomposition(`self`)`

Singular value decomposition of the graph Laplacian

---

### RandomContextualBandit.make_action_labels

`fitr.environments.make_action_labels(`self`)`

Creates labels for the actions (for plotting) if none provided

---

### RandomContextualBandit.make_digraph

`fitr.environments.make_digraph(`self`)`

Creates a `networkx` `DiGraph` object from the transition tensor for the purpose of plotting and some other analyses.

---

### RandomContextualBandit.make_state_labels

`fitr.environments.make_state_labels(`self`)`

Creates labels for the states (for plotting) if none provided

---

### RandomContextualBandit.make_undirected_graph

`fitr.environments.make_undirected_graph(`self`)`

Converts the DiGraph to undirected and computes some stats

---

### RandomContextualBandit.observation

`fitr.environments.observation(`self`)`

Samples an initial state from the start-state distribution $p(\mathbf{x})$

$$\mathbf{x}_0 \sim p(\mathbf{x})$$

Returns:

A one-hot vector `ndarray((nstates,))` indicating the starting state.

Examples:

```
x = env.observation()
```

---

**RandomContextualBandit.plot_action_outcome_probabilities**

`fitr.environments.plot_action_outcome_probabilities(self, figsize=None, outfile=None`

Plots the probabilities of different outcomes given actions.

Each plot is a heatmap for a starting state showing the transition probabilities for each action-outcome pair within that state.

---

**RandomContextualBandit.plot_graph**

`fitr.environments.plot_graph(self, figsize=None, node_size=2000, arrowsize=20, lw=1.`

Plots the directed graph of the task

---

**RandomContextualBandit.plot_spectral_properties**

`fitr.environments.plot_spectral_properties(self, figsize=None, outfile=None, outfile`

Creates a set of subplots depicting the graph Laplacian and its spectral decomposition.

---

**RandomContextualBandit.random_action**

`fitr.environments.random_action(self)`

Samples a random one-hot action vector uniformly over the action space.

Useful for testing that your environment works, without having to create an agent.

$$\mathbf{u} \sim \text{Multinomial}\left(1, \mathbf{p} = \{p_i = \frac{1}{|\mathcal{U}|}\}_{i=1}^{|\mathcal{U}|}\right)$$

Returns:

A one-hot action vector of type `ndarray((nactions,))`

Examples:

```
u = env.random_action()
```

---

**RandomContextualBandit.step**

`fitr.environments.step(self, action)`

Executes a state transition in the environment.

Arguments:

action : A one-hot vector of type `ndarray((naction,))` indicating the action selected at the current state.

Returns:

A 3-tuple representing the next state (`ndarray((noutcomes,))`), scalar reward, and whether the current step terminates a trial (`bool`).

Raises:

`RuntimeError` if `env.observation()` not called after a previous `env.step(...)` call yielded a terminal state.

---

# Chapter 4

# Agents

## `fitr.agents`

A modular way to build and test reinforcement learning agents.

There are three main submodules:

- `fitr.agents.policies`: which describe a class of functions essentially representing $f : \mathcal{X} \to \mathcal{U}$
- `fitr.agents.value_functions`: which describe a class of functions essentially representing $\mathcal{V} : \mathcal{X} \to \mathbb{R}$ and/or $\mathcal{Q} : \mathcal{Q} \times \mathcal{U} \to \mathbb{R}$
- `fitr.agents.agents`: classes of agents that are combinations of policies and value functions, along with some convenience functions for generating data from `fitr.environments.Graph` environments.

## SoftmaxPolicy

`fitr.agents.policies.SoftmaxPolicy()`

Action selection by sampling from a multinomial whose parameters are given by a softmax.

Action sampling is

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})).$$

Parameters of that distribution are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i^{|\mathbf{v}|} e^{\beta v_i}}.$$

Arguments:

- **inverse_softmax_temp**: Inverse softmax temperature $\beta$
- **rng**: `np.random.RandomState` object

---

**SoftmaxPolicy.action_prob**

`fitr.agents.policies.action_prob(self, x)`

Computes the softmax

---

**SoftmaxPolicy.log_prob**

`fitr.agents.policies.log_prob(self, x)`

Computes the log-probability of an action $\mathbf{u}$

$$\log p(\mathbf{u}|\mathbf{v}) = \beta\mathbf{v} - \log \sum_{v_i} e^{\beta\mathbf{v}_i}$$

Arguments:

- **x**: State vector of type `ndarray((nstates,))`

Returns:

Scalar log-probability

---

**SoftmaxPolicy.sample**

`fitr.agents.policies.sample(self, x)`

Samples from the action distribution

---

**StickySoftmaxPolicy**

`fitr.agents.policies.StickySoftmaxPolicy()`

Action selection by sampling from a multinomial whose parameters are given by a softmax, but with accounting for the tendency to perseverate (i.e. choosing the previously used action without considering its value).

Let $\mathbf{u}_{t-1} = (u_{t-1}^{(i)})_{i=1}^{|\mathcal{U}|}$ be a one hot vector representing the action taken at the last step, and $\beta^\rho$ be an inverse softmax temperature for the influence of this last action.

Action sampling is thus:

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v}, \mathbf{u}_{t-1})).$$

Parameters of that distribution are

$$p(\mathbf{u}|\mathbf{v}, \mathbf{u}_{t-1}) = \varsigma(\mathbf{v}, \mathbf{u}_{t-1}) = \frac{e^{\beta \mathbf{v} + \beta^\rho \mathbf{u}_{t-1}}}{\sum_i^{|\mathbf{v}|} e^{\beta v_i + \beta^\rho u_{t-1}^{(i)}}}.$$

Arguments:

- **inverse_softmax_temp**: Inverse softmax temperature $\beta$
- **perseveration**: Inverse softmax temperature $\beta^\rho$ capturing the tendency to repeat the last action taken.
- **rng**: `np.random.RandomState` object

---

**StickySoftmaxPolicy.action_prob**

`fitr.agents.policies.action_prob(self, x)`

Computes the softmax

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` vector of action probabilities

---

**StickySoftmaxPolicy.log_prob**

`fitr.agents.policies.log_prob(self, x)`

Computes the log-probability of an action $\mathbf{u}$

$$\log p(\mathbf{u}|\mathbf{v}, \mathbf{u}_{t-1}) = (\beta \mathbf{v} + \beta^\rho \mathbf{u}_{t-1}) - \log \sum_{v_i} e^{\beta \mathbf{v}_i + \beta^\rho u_{t-1}^{(i)}}$$

Arguments:

- **x**: State vector of type `ndarray((nstates,))`

Returns:

Scalar log-probability

---

**StickySoftmaxPolicy.sample**

`fitr.agents.policies.sample(self, x)`

Samples from the action distribution

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` one-hot action vector

---

## EpsilonGreedyPolicy

`fitr.agents.policies.EpsilonGreedyPolicy()`

A policy that takes the maximally valued action with probability $1 - \epsilon$, otherwise chooses randomlyself.

Arguments:

- **epsilon**: Probability of not taking the action with highest value
- **rng**: `numpy.random.RandomState` object

---

### EpsilonGreedyPolicy.action_prob

`fitr.agents.policies.action_prob(self, x)`

Creates vector of action probabilities for e-greedy policy

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` vector of action probabilities

---

### EpsilonGreedyPolicy.sample

`fitr.agents.policies.sample(self, x)`

Samples from the action distribution

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nstates,))` one-hot action vector

---

## ValueFunction

`fitr.agents.value_functions.ValueFunction()`

A general value function object.

A value function here is task specific and consists of several attributes:

- `nstates`: The number of states in the task, $|\mathcal{X}|$
- `nactions`: Number of actions in the task, $|\mathcal{U}|$
- `V`: State value function $\mathbf{v} = \mathcal{V}(\mathbf{x})$
- `Q`: State-action value function $\mathbf{Q} = \mathcal{Q}(\mathbf{x}, \mathbf{u})$
- `etrace`: An eligibility trace (optional)

Note that in general we rely on matrix-vector notation for value functions, rather than function notation. Vectors in the mathematical typesetting are by default column vectors.

Arguments:

- **env**: A `fitr.environments.Graph`

---

## ValueFunction.Qmax

`fitr.agents.value_functions.Qmax(self, x)`

Return maximal action value for given state

$$\max_{u_i} \mathcal{Q}(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

## ValueFunction.Qmean

`fitr.agents.value_functions.Qmean(self, x)`

Return mean action value for given state

$$Mean(\mathcal{Q}(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|} \mathbf{1}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### ValueFunction.Qx

`fitr.agents.value_functions.Qx(``self``, x)`

Compute action values for a given state

$$\mathcal{Q}(\mathbf{x}, :) = \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

### ValueFunction.Vx

`fitr.agents.value_functions.Vx(``self``, x)`

Compute value of state **x**

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state **x**

---

### ValueFunction.uQx

`fitr.agents.value_functions.uQx(``self``, u, x)`

Compute value of taking action **u** in state **x**

$$\mathcal{Q}(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

### ValueFunction.update

`fitr.agents.value_functions.update(`self`, x, u, r, x_, u_)`

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector
- **r**: Scalar reward
- **x_**: `ndarray((nstates,))` one-hot next-state vector
- **u_**: `ndarray((nactions,))` one-hot next-action vector

---

### DummyLearner

`fitr.agents.value_functions.DummyLearner()`

A critic/value function for the random learner

This class actually contributes nothing except identifying that a value function has been chosen for an `Agent` object

Arguments:

- **env**: A `fitr.environments.Graph`

---

### DummyLearner.Qmax

`fitr.agents.value_functions.Qmax(`self`, x)`

Return maximal action value for given state

$$\max_{u_i} \mathcal{Q}(\mathbf{x}, u_i) = \max_{\mathbf{u}'} {\mathbf{u}'}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### DummyLearner.Qmean

`fitr.agents.value_functions.Qmean(`<span style="color:blue">self</span>`, x)`

Return mean action value for given state

$$Mean(\mathcal{Q}(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|}\mathbf{1}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### DummyLearner.Qx

`fitr.agents.value_functions.Qx(`<span style="color:blue">self</span>`, x)`

Compute action values for a given state

$$\mathcal{Q}(\mathbf{x}, :) = \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

### DummyLearner.Vx

`fitr.agents.value_functions.Vx(`<span style="color:blue">self</span>`, x)`

Compute value of state **x**

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state **x**

---

### DummyLearner.uQx

`fitr.agents.value_functions.uQx(`self`, u, x)`

Compute value of taking action **u** in state **x**

$$\mathcal{Q}(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

### DummyLearner.update

`fitr.agents.value_functions.update(`self`, x, u, r, x_, u_)`

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector
- **r**: Scalar reward
- **x_**: `ndarray((nstates,))` one-hot next-state vector
- **u_**: `ndarray((nactions,))` one-hot next-action vector

---

### InstrumentalRescorlaWagnerLearner

`fitr.agents.value_functions.InstrumentalRescorlaWagnerLearner()`

Learns an instrumental control policy through one-step error-driven updates of the state-action value function

The instrumental Rescorla-Wagner rule is as follows:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r - \mathbf{u}^\top \mathbf{Q}\mathbf{x})\mathbf{u}\mathbf{x}^\top,$$

where $0 < \alpha < 1$ is the learning rate, and where the reward prediction error (RPE) is $\delta = (r - \mathbf{u}^\top \mathbf{Q}\mathbf{x})$.

$$

Arguments:

- **env**: A `fitr.environments.Graph`
- **learning_rate**: Learning rate $\alpha$

------

### InstrumentalRescorlaWagnerLearner.Qmax

`fitr.agents.value_functions.Qmax(self, x)`

Return maximal action value for given state

$$\max_{u_i} \mathcal{Q}(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

------

### InstrumentalRescorlaWagnerLearner.Qmean

`fitr.agents.value_functions.Qmean(self, x)`

Return mean action value for given state

$$Mean(\mathcal{Q}(\mathbf{x},:)) = \frac{1}{|\mathcal{U}|}\mathbf{1}^\top \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

------

**InstrumentalRescorlaWagnerLearner.Qx**

`fitr.agents.value_functions.Qx(`<span style="color:blue">self</span>`, x)`

Compute action values for a given state

$$\mathcal{Q}(\mathbf{x}, :) = \mathbf{Qx}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

**InstrumentalRescorlaWagnerLearner.Vx**

`fitr.agents.value_functions.Vx(`<span style="color:blue">self</span>`, x)`

Compute value of state **x**

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state **x**

---

**InstrumentalRescorlaWagnerLearner.uQx**

`fitr.agents.value_functions.uQx(`<span style="color:blue">self</span>`, u, x)`

Compute value of taking action **u** in state **x**

$$\mathcal{Q}(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Qx}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

**InstrumentalRescorlaWagnerLearner.update**

`fitr.agents.value_functions.update(`<span style="color:blue">`self`</span>`, x, u, r, x_, u_)`

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector
- **r**: Scalar reward
- **x_**: `ndarray((nstates,))` one-hot next-state vector
- **u_**: `ndarray((nactions,))` one-hot next-action vector

---

## QLearner

`fitr.agents.value_functions.QLearner()`

Learns an instrumental control policy through Q-learning

The Q-learning rule is as follows:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^{\top}\mathbf{Q}\mathbf{x}' - \mathbf{u}^{\top}\mathbf{Q}\mathbf{x})\mathbf{z},$$

where $0 < \alpha < 1$ is the learning rate, $0 \leq \gamma \leq 1$ is a discount factor, and where the reward prediction error (RPE) is $\delta = (r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^{\top}\mathbf{Q}\mathbf{x}' - \mathbf{u}^{\top}\mathbf{Q}\mathbf{x})$. We have also included an eligibility trace $\mathbf{z}$ defined as

$$\mathbf{z} = \mathbf{u}\mathbf{x}^{\top} + \gamma\lambda\mathbf{z}$$

Arguments:

- **env**: A `fitr.environments.Graph`
- **learning_rate**: Learning rate $\alpha$
- **discount_factor**: Discount factor $\gamma$
- **trace_decay**: Eligibility trace decay $\lambda$

---

## QLearner.Qmax

`fitr.agents.value_functions.Qmax(`<span style="color:blue">`self`</span>`, x)`

Return maximal action value for given state

$$\max_{u_i} \mathcal{Q}(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^{\top}\mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### QLearner.Qmean

`fitr.agents.value_functions.Qmean(self, x)`

Return mean action value for given state

$$Mean(\mathcal{Q}(\mathbf{x},:)) = \frac{1}{|\mathcal{U}|}\mathbf{1}^\top\mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### QLearner.Qx

`fitr.agents.value_functions.Qx(self, x)`

Compute action values for a given state

$$\mathcal{Q}(\mathbf{x},:) = \mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

### QLearner.Vx

`fitr.agents.value_functions.Vx(self, x)`

Compute value of state $\mathbf{x}$

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state $\mathbf{x}$

---

### QLearner.uQx

`fitr.agents.value_functions.uQx(`<span style="color:blue">self</span>`, u, x)`

Compute value of taking action $\mathbf{u}$ in state $\mathbf{x}$

$$\mathcal{Q}(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Q} \mathbf{x}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action $\mathbf{u}$ in state $\mathbf{x}$

---

### QLearner.update

`fitr.agents.value_functions.update(`<span style="color:blue">self</span>`, x, u, r, x_, u_)`

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector
- **r**: Scalar reward
- **x_**: `ndarray((nstates,))` one-hot next-state vector
- **u_**: `ndarray((nactions,))` one-hot next-action vector

---

### SARSALearner

`fitr.agents.value_functions.SARSALearner()`

Learns an instrumental control policy through the SARSA learning rule

The SARSA learning rule is as follows:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r + \gamma \mathbf{u}'^{\top}\mathbf{Q}\mathbf{x}' - \mathbf{u}^{\top}\mathbf{Q}\mathbf{x})\mathbf{z},$$

where $0 < \alpha < 1$ is the learning rate, $0 \leq \gamma \leq 1$ is a discount factor, and where the reward prediction error (RPE) is $\delta = (r + \gamma \mathbf{u}'^{\top}\mathbf{Q}\mathbf{x}' - \mathbf{u}^{\top}\mathbf{Q}\mathbf{x})$. We have also included an eligibility trace $\mathbf{z}$ defined as

$$\mathbf{z} = \mathbf{u}\mathbf{x}^{\top} + \gamma\lambda\mathbf{z}$$

Arguments:

- **env**: A `fitr.environments.Graph`
- **learning_rate**: Learning rate $\alpha$
- **discount_factor**: Discount factor $\gamma$
- **trace_decay**: Eligibility trace decay $\lambda$

---

### SARSALearner.Qmax

`fitr.agents.value_functions.Qmax(self, x)`

Return maximal action value for given state

$$\max_{u_i} \mathcal{Q}(\mathbf{x}, u_i) = \max_{\mathbf{u}'} \mathbf{u}'^{\top}\mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### SARSALearner.Qmean

`fitr.agents.value_functions.Qmean(self, x)`

Return mean action value for given state

$$Mean(\mathcal{Q}(\mathbf{x}, :)) = \frac{1}{|\mathcal{U}|}\mathbf{1}^{\top}\mathbf{Q}\mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of the maximal action value at the given state

---

### SARSALearner.Qx

`fitr.agents.value_functions.Qx(`self`, x)`

Compute action values for a given state

$$\mathcal{Q}(\mathbf{x},:) = \mathbf{Qx}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` vector of values for actions in the given state

---

### SARSALearner.Vx

`fitr.agents.value_functions.Vx(`self`, x)`

Compute value of state **x**

$$\mathcal{V}(\mathbf{x}) = \mathbf{v}^\top \mathbf{x}$$

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of state **x**

---

### SARSALearner.uQx

`fitr.agents.value_functions.uQx(`self`, u, x)`

Compute value of taking action **u** in state **x**

$$\mathcal{Q}(\mathbf{x}, \mathbf{u}) = \mathbf{u}^\top \mathbf{Qx}$$

Arguments:

- **u**: `ndarray((nactions,))` one-hot action vector
- **x**: `ndarray((nstates,))` one-hot state vector

Returns:

Scalar value of action **u** in state **x**

---

### SARSALearner.update

`fitr.agents.value_functions.update(`<span style="color:blue">`self`</span>`, x, u, r, x_, u_)`

Updates the value function

In the context of the base `ValueFunction` class, this is merely a placeholder. The specific update rule will depend on the specific value function desired.

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector
- **r**: Scalar reward
- **x_**: `ndarray((nstates,))` one-hot next-state vector
- **u_**: `ndarray((nactions,))` one-hot next-action vector

---

### Agent

`fitr.agents.agents.Agent()`

Base class for synthetic RL agents.

Arguments:

meta : List of metadata of arbitrary type. e.g. labels, covariates, etc. params : List of parameters for the agent. Should be filled for specific agent.

---

### Agent.action

`fitr.agents.agents.action(`<span style="color:blue">`self`</span>`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

### Agent.learning

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

### Agent.reset_trace

`fitr.agents.agents.reset_trace(`self`, x, u=`None`)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

## BanditAgent

`fitr.agents.agents.BanditAgent()`

A base class for agents in bandit tasks (i.e. with one step).

Arguments:

- **task**: `fitr.environments.Graph`

---

### BanditAgent.action

`fitr.agents.agents.action(`self`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

**BanditAgent.generate_data**

`fitr.agents.agents.generate_data(`self`, ntrials)`

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

**BanditAgent.learning**

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

**BanditAgent.log_prob**

`fitr.agents.agents.log_prob(`self`, state)`

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` log-likelihood vector

---

### BanditAgent.reset_trace

`fitr.agents.agents.reset_trace(`<span style="color:blue">`self`</span>`, x, u=`<span style="color:blue">`None`</span>`)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

### MDPAgent

`fitr.agents.agents.MDPAgent()`

A base class for agents that operate on MDPs.

This mainly has implications for generating data.

Arguments:

- **task**: `fitr.environments.Graph`

---

### MDPAgent.action

`fitr.agents.agents.action(`<span style="color:blue">`self`</span>`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

### MDPAgent.generate_data

`fitr.agents.agents.generate_data(`<span style="color:blue">`self`</span>`, ntrials)`

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

## MDPAgent.learning

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

## MDPAgent.reset_trace

`fitr.agents.agents.reset_trace(`self`, x, u=`None`)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

## RandomBanditAgent

`fitr.agents.agents.RandomBanditAgent()`

An agent that simply selects random actions at each trial

---

## RandomBanditAgent.action

`fitr.agents.agents.action(`self`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

**RandomBanditAgent.generate_data**

`fitr.agents.agents.generate_data(`self`, ntrials)`

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

**RandomBanditAgent.learning**

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

**RandomBanditAgent.log_prob**

`fitr.agents.agents.log_prob(`self`, state)`

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` log-likelihood vector

---

**RandomBanditAgent.reset_trace**

`fitr.agents.agents.reset_trace(`self`, x, u=`None`)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

## RandomMDPAgent

`fitr.agents.agents.RandomMDPAgent()`

An agent that simply selects random actions at each trial

### Notes

This has been specified as an `OnPolicyAgent` arbitrarily.

---

**RandomMDPAgent.action**

`fitr.agents.agents.action(`self`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

**RandomMDPAgent.generate_data**

`fitr.agents.agents.generate_data(`self`, ntrials)`

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

**RandomMDPAgent.learning**

`fitr.agents.agents.learning(`<code>self</code>`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

**RandomMDPAgent.reset_trace**

`fitr.agents.agents.reset_trace(`<code>self</code>`, x, u=None)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

**SARSASoftmaxAgent**

`fitr.agents.agents.SARSASoftmaxAgent()`

An agent that uses the SARSA learning rule and a softmax policy

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i^{|\mathbf{v}|} e^{\beta v_i}}.$$

The value function is SARSA:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r + \gamma \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x})\mathbf{z},$$

where $0 < \alpha < 1$ is the learning rate, $0 \leq \gamma \leq 1$ is a discount factor, and where the reward prediction error (RPE) is $\delta = (r + \gamma \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$. We have also included an eligibility trace $\mathbf{z}$ defined as

$$\mathbf{z} = \mathbf{u}\mathbf{x}^\top + \gamma\lambda\mathbf{z}$$

Arguments:

- **task**: `fitr.environments.Graph`
- **learning_rate**: Learning rate $\alpha$
- **discount_factor**: Discount factor $\gamma$
- **trace_decay**: Eligibility trace decay $\lambda$
- **inverse_softmax_temp**: Inverse softmax temperature $\beta$
- **rng**: `np.random.RandomState`

---

## SARSASoftmaxAgent.action

`fitr.agents.agents.action(`self`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

## SARSASoftmaxAgent.generate_data

`fitr.agents.agents.generate_data(`self`, ntrials)`

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

## SARSASoftmaxAgent.learning

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

### SARSASoftmaxAgent.reset_trace

`fitr.agents.agents.reset_trace(self, x, u=None)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

## QLearningSoftmaxAgent

`fitr.agents.agents.QLearningSoftmaxAgent()`

An agent that uses the Q-learning rule and a softmax policy

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i^{|\mathbf{v}|} e^{\beta v_i}}.$$

The value function is Q-learning:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x}) \mathbf{z},$$

where $0 < \alpha < 1$ is the learning rate, $0 \leq \gamma \leq 1$ is a discount factor, and where the reward prediction error (RPE) is $\delta = (r + \gamma \max_{\mathbf{u}'} \mathbf{u}'^\top \mathbf{Q} \mathbf{x}' - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$. The eligibility trace $\mathbf{z}$ is defined as

$$\mathbf{z} = \mathbf{u} \mathbf{x}^\top + \gamma \lambda \mathbf{z}$$

Arguments:

- **task**: `fitr.environments.Graph`
- **learning_rate**: Learning rate $\alpha$
- **discount_factor**: Discount factor $\gamma$

- **trace_decay**: Eligibility trace decay $\lambda$
- **inverse_softmax_temp**: Inverse softmax temperature $\beta$
- **rng**: `np.random.RandomState`

---

### QLearningSoftmaxAgent.action

`fitr.agents.agents.action(`self`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

### QLearningSoftmaxAgent.generate_data

`fitr.agents.agents.generate_data(`self`, ntrials)`

For the parent agent, this function generates data from a Markov Decision Process (MDP) task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

### QLearningSoftmaxAgent.learning

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

**QLearningSoftmaxAgent.reset_trace**

`fitr.agents.agents.reset_trace(`<span style="color:blue">`self`</span>`, x, u=`<span style="color:blue">`None`</span>`)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

## RWSoftmaxAgent

`fitr.agents.agents.RWSoftmaxAgent()`

An instrumental Rescorla-Wagner agent with a softmax policy

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i^{|\mathbf{v}|} e^{\beta v_i}}.$$

The value function is the Rescorla-Wagner learning rule:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(r - \mathbf{u}^\top \mathbf{Q} \mathbf{x})\mathbf{u}\mathbf{x}^\top,$$

where $0 < \alpha < 1$ is the learning rate, $0 \leq \gamma \leq 1$ is a discount factor, and where the reward prediction error (RPE) is $\delta = (r - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$.

Arguments:

- **task**: `fitr.environments.Graph`
- **learning_rate**: Learning rate $\alpha$
- **inverse_softmax_temp**: Inverse softmax temperature $\beta$
- **rng**: `np.random.RandomState`

---

**RWSoftmaxAgent.action**

`fitr.agents.agents.action(`<span style="color:blue">`self`</span>`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

### RWSoftmaxAgent.generate_data

`fitr.agents.agents.generate_data(`self`, ntrials)`

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

### RWSoftmaxAgent.learning

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

### RWSoftmaxAgent.log_prob

`fitr.agents.agents.log_prob(`self`, state)`

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` log-likelihood vector

---

**RWSoftmaxAgent.reset_trace**

`fitr.agents.agents.reset_trace(`self`, x, u=`None`)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

**RWSoftmaxAgentRewardSensitivity**

`fitr.agents.agents.RWSoftmaxAgentRewardSensitivity()`

An instrumental Rescorla-Wagner agent with a softmax policy, whose experienced reward is scaled by a factor $\rho$.

The softmax policy selects actions from a multinomial

$$\mathbf{u} \sim \text{Multinomial}(1, \mathbf{p} = \varsigma(\mathbf{v})),$$

whose parameters are

$$p(\mathbf{u}|\mathbf{v}) = \varsigma(\mathbf{v}) = \frac{e^{\beta \mathbf{v}}}{\sum_i^{|\mathbf{v}|} e^{\beta v_i}}.$$

The value function is the Rescorla-Wagner learning rule with scaled reward $\rho r$:

$$\mathbf{Q} \leftarrow \mathbf{Q} + \alpha(\rho r - \mathbf{u}^\top \mathbf{Q} \mathbf{x})\mathbf{u}\mathbf{x}^\top,$$

where $0 < \alpha < 1$ is the learning rate, $0 \leq \gamma \leq 1$ is a discount factor, and where the reward prediction error (RPE) is $\delta = (\rho r - \mathbf{u}^\top \mathbf{Q} \mathbf{x})$.

Arguments:

- **task**: `fitr.environments.Graph`
- **learning_rate**: Learning rate $\alpha$
- **inverse_softmax_temp**: Inverse softmax temperature $\beta$
- **reward_sensitivity**: Reward sensitivity parameter $\rho$
- **rng**: `np.random.RandomState`

---

**RWSoftmaxAgentRewardSensitivity.action**

`fitr.agents.agents.action(`self`, state)`

Selects an action given the current state of environment.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

---

### RWSoftmaxAgentRewardSensitivity.generate_data

`fitr.agents.agents.generate_data(`self`, ntrials)`

For the parent agent, this function generates data from a bandit task

Arguments:

- **ntrials**: `int` number of trials

Returns:

`fitr.data.BehaviouralData`

---

### RWSoftmaxAgentRewardSensitivity.learning

`fitr.agents.agents.learning(`self`, state, action, reward, next_state, next_action)`

Updates the model's parameters.

The implementation will vary depending on the type of agent and environment.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector
- **action**: `ndarray((nactions,))` one-hot action vector
- **reward**: scalar reward
- **next_state**: `ndarray((nstates,))` one-hot next-state vector
- **next_action**: `ndarray((nactions,))` one-hot action vector

---

### RWSoftmaxAgentRewardSensitivity.log_prob

`fitr.agents.agents.log_prob(`self`, state)`

Computes the log-likelihood over actions for a given state under the present agent parameters.

Presently this only works for the state-action value function. In all other cases, you should define your own log-likelihood function. However, this can be used as a template.

Arguments:

- **state**: `ndarray((nstates,))` one-hot state vector

Returns:

`ndarray((nactions,))` log-likelihood vector

---

### RWSoftmaxAgentRewardSensitivity.reset_trace

`fitr.agents.agents.reset_trace(`<span style="color:blue">`self`</span>`, x, u=`<span style="color:blue">`None`</span>`)`

For agents with eligibility traces, this resets the eligibility trace (for episodic tasks)

Arguments:

- **x**: `ndarray((nstates,))` one-hot state vector
- **u**: `ndarray((nactions,))` one-hot action vector (optional)

---

# Chapter 5

# Data

## **fitr.data**

A module containing a generic class for behavioural data.

### BehaviouralData

`fitr.data.BehaviouralData()`

A flexible and generic object to store and process behavioural data across tasks

Arguments:

- **ngroups**: Integer number of groups represented in the dataset. Only > 1 if data are merged
- **nsubjects**: Integer number of subjects in dataset
- **ntrials**: Integer number of trials done by each subject
- **dict**: Dictionary storage indexed by subject.
- **params**: `ndarray((nsubjects, nparams + 1))` parameters for each (simulated) subject
- **meta**: Array of covariates of type `ndarray((nsubjects, nmetadata_features+1))`
- **tensor**: Tensor representation of the behavioural data of type `ndarray((nsubjects, ntrials, nfeatures))`

---

### BehaviouralData.add_subject

`fitr.data.add_subject(self, subject_index, parameters, subject_meta)`

Appends a new subject to the dataset

Arguments:

- **subject_index**: Integer identification for subject
- **parameters**: `list` of parameters for the subject
- **subject_meta**: Some covariates for the subject (`list`)

---

**BehaviouralData.initialize_data_dictionary**

```
fitr.data.initialize_data_dictionary(self)
```

---

**BehaviouralData.make_behavioural_ngrams**

```
fitr.data.make_behavioural_ngrams(self, n)
```

Creates N-grams of behavioural data

---

**BehaviouralData.make_cooccurrence_matrix**

```
fitr.data.make_cooccurrence_matrix(self, k, dtype=<class 'numpy.float32'>)
```

---

**BehaviouralData.make_tensor_representations**

```
fitr.data.make_tensor_representations(self)
```

Creates a tensor with all subjects' data

**Notes**

Assumes that all subjects did same number of trials.

---

**BehaviouralData.numpy_tensor_to_bdf**

```
fitr.data.numpy_tensor_to_bdf(self, X)
```

Creates `BehaviouralData` formatted set from a dataset stored in a numpy `ndarray`.

Arguments:

- **X**: `ndarray((nsubjects, ntrials, m))` with `m` being the size of flattened single-trial data

---

**BehaviouralData.unpack_tensor**

```
fitr.data.unpack_tensor(self, x_dim, u_dim, r_dim=1, terminal_dim=1, get='sarsat')
```

Unpacks data stored in tensor format into separate arrays for states, actions, rewards, next states, and next actions.

Arguments:

x_dim : Task state space dimensionality (`int`) u_dim : Task action space dimensionality (`int`) r_dim : Reward dimensionality (`int`, default=1) terminal_dim : Dimensionality of the terminal state indicator (`int` , default=1) get : String indicating the order that data are stored in the array. Can also be shortened such that fewer elements are returned. For example, the default is `sarsat`.

Returns:

List with data, where each element is in the order of the argument `get`

---

### BehaviouralData.update

`fitr.data.update(`self`, subject_index, behav_data)`

Adds behavioural data to the dataset

Arguments:

- **subject_index**: Integer index for the subject
- **behav_data**: 1-dimensional `ndarray` of flattened data

---

### merge_behavioural_data

`fitr.data.merge_behavioural_data(datalist)`

Combines BehaviouralData objects.

Arguments:

- **datalist**: List of BehaviouralData objects

Returns:

`BehaviouralData` with data from multiple groups merged.

---

# Chapter 6

# Inference

## `fitr.inference`

Methods for inferring the parameters of generative models for reinforcement learning data.

### OptimizationResult

`fitr.inference.optimization_result.OptimizationResult()`

Container for the results of an optimization run on a generative model of behavioural data

Arguments:

- **subject_id**: `ndarray((nsubjects,))` or `None` (default). Integer ids for subjects
- **xmin**: `ndarray((nsubjects,nparams))` or `None` (default). Parameters that minimize objective function
- **fmin**: `ndarray((nsubjects,))` or `None` (default). Value of objective function at minimum
- **fevals**: `ndarray((nsubjects,))` or `None` (default). Number of function evaluations required to minimize objective function
- **niters**: `ndarray((nsubjects,))` or `None` (default). Number of iterations required to minimize objective function
- **lme**: `ndarray((nsubjects,))` or `None` (default). Log model evidence
- **bic**: `ndarray((nsubjects,))` or `None` (default). Bayesian Information Criterion
- **hess_inv**: `ndarray((nsubjects,nparams,nparams))` or `None` (default). Inverse Hessian at the optimum.
- **err**: `ndarray((nsubjects,nparams))` or `None` (default). Error of estimates at optimum.

---

### OptimizationResult.transform_xmin

`fitr.inference.optimization_result.transform_xmin(self, transforms, inplace=False)`

Rescales the parameter estimates.

Arguments:

- **transforms**: `list`. Transformation functions where `len(transforms) == self.xmin.shape[1]`
- **inplace**: `bool`. Whether to change the values in `self.xmin`. Default is `False`, which returns an `ndarray((nsubjects, nparams))` of the transformed parameters.

Returns:

`ndarray((nsubjects, nparams))` of the transformed parameters if `inplace=False`

---

## mlepar

`fitr.inference.mle_parallel.mlepar(f, data, nparams, minstarts=2, maxstarts=10, init`

Computes maximum likelihood estimates using parallel CPU resources.

Wraps over the `fitr.optimization.mle_parallel.mle` function.

Arguments:

- **f**: Likelihood function
- **data**: A subscriptable object whose first dimension indexes subjects
- **optimizer**: Optimization function (currently only `l_bfgs_b` supported)
- **nparams**: `int` number of parameters to be estimated
- **minstarts**: `int`. Minimum number of restarts with new initial values
- **maxstarts**: `int`. Maximum number of restarts with new initial values
- **init_sd**: Standard deviation for Gaussian initial values

Returns:

`fitr.inference.OptimizationResult`

---

## l_bfgs_b

`fitr.inference.mle_parallel.l_bfgs_b(f, i, data, nparams, minstarts=2, maxstarts=10,`

Minimizes the negative log-probability of data with respect to some parameters under function `f` using the L-BFGS-B algorithm.

This function is specified for use with parallel CPU resources.

Arguments:

- **f**: Log likelihood function
- **i**: `int`. Subject being optimized (slices first dimension of `data`)
- **data**: Object subscriptable along first dimension to indicate subject being optimized
- **nparams**: `int`. Number of parameters in the model
- **minstarts**: `int`. Minimum number of restarts with new initial values
- **maxstarts**: `int`. Maximum number of restarts with new initial values
- **init_sd**: Standard deviation for Gaussian initial values

Returns:

- **i**: `int`. Subject being optimized (slices first dimension of `data`)
- **xmin**: `ndarray((nparams,))`. Parameter values at optimum
- **fmin**: Scalar objective function value at optimum
- **fevals**: `int`. Number of function evaluations
- **niters**: `int`. Number of iterations
- **lme_**: Scalar log-model evidence at optimum
- **bic_**: Scalar Bayesian Information Criterion at optimum
- **hess_inv**: `ndarray((nparams, nparams))`. Inv at optimum

---

# Chapter 7

# Criticism

## **`fitr.criticism`**

Methods for criticism of model fits.

### actual_estimate

`fitr.criticism.plotting.actual_estimate(y_true, y_pred, xlabel='Actual', ylabel='Est`

Plots parameter estimates against the ground truth values.

Arguments:

- **y_true**: `ndarray(nsamples)`. Vector of ground truth parameters
- **y_pred**: `ndarray(nsamples)`. Vector of parameter estimates
- **xlabel**: `str`. Label for x-axis
- **ylabel**: `str`. Label for y-axis
- **corr**: `bool`. Whether to plot correlation coefficient.
- **figsize**: `tuple`. Figure size (inches).

Returns:

`matplotlib.pyplot.Figure`

---

# Chapter 8

# Metrics

## **fitr.metrics**

Metrics and performance statistics.

### bic

`fitr.metrics.bic(log_prob, nparams, ntrials)`

Bayesian Information Criterion (BIC)

Arguments:

- **log_prob**: Log probability
- **nparams**: Number of parameters in the model
- **ntrials**: Number of trials in the time series

Returns:

Scalar estimate of BIC.

---

### linear_correlation

`fitr.metrics.linear_correlation(X, Y)`

Linear correlation coefficient.

Will compute the following formula

$$\rho = \frac{\mathbf{x}^\top \mathbf{y}}{\|\mathbf{x}Vert \cdot \|\mathbf{y}Vert}$$

where each vector $\mathbf{x}$ and $\mathbf{y}$ are rows of the matrices $\mathbf{X}$ and $\mathbf{Y}$, respectively.

Arguments:

- **X**: `ndarray((nsamples, nfeatures))` of dimension 1 or 2. If `X` is a 1D array, it will be converted to 2D prior to computation
- **Y**: `ndarray((nsamples, nfeatures))` of dimension 1 or 2. If `Y` is a 1D array, it will be converted to 2D prior to computation

Returns:

- **rho**: `ndarray((nfeatures,))`. Correlation coefficient(s)

TODO:

- [ ] Create error raised when X and Y are not same dimension

---

### lme

`fitr.metrics.lme(log_prob, nparams, hess_inv)`

Laplace approximation to the log model evidence

Arguments:

- **log_prob**: Log probability
- **nparams**: Number of parameters in the model
- **hess_inv**: Hessian at the optimum (shape is $K \times K$)

Returns:

Scalar approximation of the log model evidence

---

### log_loss

`fitr.metrics.log_loss(p, q)`

Computes log loss.

$$\mathcal{L} = \mathbf{p}^\top \log \mathbf{q} + (1 - \mathbf{p})^\top \log(1 - \mathbf{q})$$

Arguments:

- **p**: Binary vector of true labels `ndarray((nsamples,))`
- **q**: Vector of estimates (between 0 and 1) of type `ndarray((nsamples,))`

Returns:

Scalar log loss

---

# Chapter 9

# Utilities

## **fitr.utils**

Functions used across `fitr`.

### logsumexp

`fitr.utils.logsumexp(x)`

Numerically stable logsumexp.

Computed as follows:

$$\max x + \log \sum_x e^{x - \max x}$$

Arguments:

- **x**: 'ndarray(shape=(nactions,))''

Returns:

`float`

---

### relu

`fitr.utils.relu(x, a_max=None)`

Rectified linearity

$$\mathbf{x}' = \max(x_i, 0)_{i=1}^{|\mathbf{x}|}$$

Arguments:

- **x**: Vector of inputs

   • **a_max**: Upper bound at which to clip values of x

Returns:

Exponentiated values of x.

---

## scale_data

`fitr.utils.scale_data(X, axis=0, with_mean=True, with_var=True)`

Rescales data by subtracting mean and dividing by variance

$$\mathbf{x}' = \frac{\mathbf{x} - \frac{1}{n}\mathbf{1}^\top\mathbf{x}}{Var(\mathbf{x})}$$

Arguments:

   • **X**: `ndarray((nsamples, [nfeatures]))`. Data. May be 1D or 2D.
   • **with_mean**: `bool`. Whether to subtract the mean
   • **with_var**: `bool`. Whether to divide by variance

Returns:

`ndarray(X.shape)`. Rescaled data.

---

## sigmoid

`fitr.utils.sigmoid(x, a_min=-10, a_max=10)`

Sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Arguments:

   • **x**: Vector
   • **a_min**: Lower bound at which to clip values of x
   • **a_max**: Upper bound at which to clip values of x

Returns:

Vector between 0 and 1 of size `x.shape`

---

**softmax**

`fitr.utils.softmax(x)`

Computes the softmax function

$$p(\mathbf{x}) = \frac{e^{\mathbf{x} - \max_i x_i}}{\mathbf{1}^\top e^{\mathbf{x} - \max_i x_i}}$$

Arguments:

- **x**: Softmax logits (`ndarray((N,))`)

Returns:

Vector of probabilities of size `ndarray((N,))`

---

**stable_exp**

`fitr.utils.stable_exp(x, a_min=-10, a_max=10)`

Clipped exponential function

Avoids overflow by clipping input values.

Arguments:

- **x**: Vector of inputs
- **a_min**: Lower bound at which to clip values of `x`
- **a_max**: Upper bound at which to clip values of `x`

Returns:

Exponentiated values of `x`.

---